

# PYTHON CRASH COURSE

A HANDS-ON, PROJECT-BASED  
INTRODUCTION TO PROGRAMMING

ERIC MATTHES



WOW! eBook  
[www.wowebook.org](http://www.wowebook.org)



# PROJECT 3

## WEB APPLICATIONS



# 18

## GETTING STARTED WITH DJANGO



Behind the scenes, today's websites are actually rich applications that act like fully developed desktop applications. Python has a great set of tools for building web applications. In this chapter you'll learn how to use Django (<http://djangoproject.com/>) to build a project called Learning Log—an online journal system that lets you keep track of information you've learned about particular topics.

We'll write a specification for this project, and then we'll define models for the data the app will work with. We'll use Django's admin system to enter some initial data and then learn to write views and templates so Django can build the pages of our site.

Django is a *web framework*—a set of tools designed to help you build interactive websites. Django can respond to page requests and make it

easier to read and write to a database, manage users, and much more. In Chapters 19 and 20 we'll refine the Learning Log project and then deploy it to a live server so you (and your friends) can use it.

## Setting Up a Project

When beginning a project, you first need to describe the project in a specification, or *spec*. Then you'll set up a virtual environment to build the project in.

### Writing a Spec

A full spec details the project goals, describes the project's functionality, and discusses its appearance and user interface. Like any good project or business plan, a spec should keep you focused and help keep your project on track. We won't write a full project spec here, but we'll lay out a few clear goals to keep our development process focused. Here's the spec we'll use:

We'll write a web app called Learning Log that allows users to log the topics they're interested in and to make journal entries as they learn about each topic. The Learning Log home page should describe the site and invite users to either register or log in. Once logged in, a user should be able to create new topics, add new entries, and read and edit existing entries.

When you learn about a new topic, keeping a journal of what you've learned can be helpful in tracking and revisiting information. A good app makes this process efficient.

### Creating a Virtual Environment

To work with Django, we'll first set up a virtual environment to work in. A *virtual environment* is a place on your system where you can install packages and isolate them from all other Python packages. Separating one project's libraries from other projects is beneficial and will be necessary when we deploy Learning Log to a server in Chapter 20.

Create a new directory for your project called *learning\_log*, switch to that directory in a terminal, and create a virtual environment. If you're using Python 3, you should be able to create a virtual environment with the following command:

---

```
learning_log$ python -m venv ll_env
learning_log$
```

---

Here we're running the `venv` module and using it to create a virtual environment named *ll\_env*. If this works, move on to "Activating the Virtual Environment" on page 399. If it doesn't work, read the next section, "Installing virtualenv."

## Installing virtualenv

If you're using an earlier version of Python or if your system isn't set up to use the `venv` module correctly, you can install the `virtualenv` package. To install `virtualenv`, enter the following:

---

```
$ pip install --user virtualenv
```

---

Keep in mind that you might need to use a slightly different version of this command. (If you haven't used `pip` yet, see "Installing Python Packages with `pip`" on page 237.)

### NOTE

*If you're using Linux and this still doesn't work, you can install `virtualenv` through your system's package manager. On Ubuntu, for example, the command `sudo apt-get install python-virtualenv` will install `virtualenv`.*

Change to the `learning_log` directory in a terminal, and create a virtual environment like this:

---

```
learning_log$ virtualenv ll_env
New python executable in ll_env/bin/python
Installing setuptools, pip...done.
learning_log$
```

---

### NOTE

*If you have more than one version of Python installed on your system, you should specify the version for `virtualenv` to use. For example, the command `virtualenv ll_env --python=python3` will create a virtual environment that uses Python 3.*

## Activating the Virtual Environment

Now that we have a virtual environment set up, we need to activate it with the following command:

---

```
learning_log$ source ll_env/bin/activate
❶ (ll_env)learning_log$
```

---

This command runs the script `activate` in `ll_env/bin`. When the environment is active, you'll see the name of the environment in parentheses, as shown at ❶; then you can install packages to the environment and use packages that have already been installed. Packages you install in `ll_env` will be available only while the environment is active.

### NOTE

*If you're using Windows, use the command `ll_env\Scripts\activate` (without the word `source`) to activate the virtual environment.*

To stop using a virtual environment, enter `deactivate`:

---

```
(ll_env)learning_log$ deactivate
learning_log$
```

---

The environment will also become inactive if you close the terminal it's running in.

## Installing Django

Once you've created your virtual environment and activated it, install Django:

---

```
(ll_env)learning_log$ pip install Django
Installing collected packages: Django
Successfully installed Django
Cleaning up...
(ll_env)learning_log$
```

---

Because we're working in a virtual environment, this command is the same on all systems. There's no need to use the `--user` flag, and there's no need to use longer commands like `python -m pip install package_name`.

Keep in mind that Django will be available only when the environment is active.

## Creating a Project in Django

Without leaving the active virtual environment (remember to look for `ll_env` in parentheses), enter the following commands to create a new project:

---

```
❶ (ll_env)learning_log$ django-admin.py startproject learning_log .
❷ (ll_env)learning_log$ ls
learning_log ll_env manage.py
❸ (ll_env)learning_log$ ls learning_log
__init__.py settings.py urls.py wsgi.py
```

---

The command at ❶ tells Django to set up a new project called *learning\_log*. The dot at the end of the command creates the new project with a directory structure that will make it easy to deploy the app to a server when we're finished developing it.

### NOTE

*Don't forget this dot, or you may run into some configuration issues when we deploy the app. If you forget the dot, delete the files and folders that were created (except `ll_env`), and run the command again.*

Running the `ls` command (dir on Windows) ❷ shows that Django has created a new directory called *learning\_log*. It also created a file called *manage.py*, which is a short program that takes in commands and feeds them to the relevant part of Django to run them. We'll use these commands to manage tasks like working with databases and running servers.

The *learning\_log* directory contains four files ❸, the most important of which are *settings.py*, *urls.py*, and *wsgi.py*. The *settings.py* file controls how Django interacts with your system and manages your project. We'll modify a few of these settings and add some settings of our own as the project

evolves. The *urls.py* file tells Django which pages to build in response to browser requests. The *wsgi.py* file helps Django serve the files it creates. The filename is an acronym for *web server gateway interface*.

## Creating the Database

Because Django stores most of the information related to a project in a database, we need to create a database that Django can work with. To create the database for the Learning Log project, enter the following command (still in an active environment):

---

```
(ll_env)learning_log$ python manage.py migrate
❶ Operations to perform:
  Synchronize unmigrated apps: messages, staticfiles
  Apply all migrations: contenttypes, sessions, auth, admin
  --snip--
  Applying sessions.0001_initial... OK
❷ (ll_env)learning_log$ ls
db.sqlite3  learning_log  ll_env  manage.py
```

---

Any time we modify a database, we say we're *migrating* the database. Issuing the migrate command for the first time tells Django to make sure the database matches the current state of the project. The first time we run this command in a new project using SQLite (more about SQLite in a moment), Django will create a new database for us. At ❶ Django reports that it will make the database tables needed to store the information we'll use in this project (*Synchronize unmigrated apps*), and then make sure the database structure matches the current code (*Apply all migrations*).

Running the ls command shows that Django created another file called *db.sqlite3* ❷. SQLite is a database that runs off a single file; it's ideal for writing simple apps because you won't have to pay much attention to managing the database.

## Viewing the Project

Let's make sure that Django has set up the project properly. Enter the runserver command as follows:

---

```
(ll_env)learning_log$ python manage.py runserver
Performing system checks...

❶ System check identified no issues (0 silenced).
  July 15, 2015 - 06:23:51
❷ Django version 1.8.4, using settings 'learning_log.settings'
❸ Starting development server at http://127.0.0.1:8000/
  Quit the server with CONTROL-C.
```

---

Django starts a server so you can view the project on your system to see how well it works. When you request a page by entering a URL in a browser, the Django server responds to that request by building the appropriate page and sending that page to the browser.



At ❶ Django checks to make sure the project is set up properly; at ❷ it reports the version of Django in use and the name of the settings file being used; and at ❸ it reports the URL where the project is being served. The URL `http://127.0.0.1:8000/` indicates that the project is listening for requests on port 8000 on your computer—called a *localhost*. The term *localhost* refers to a server that only processes requests on your system; it doesn't allow anyone else to see the pages you're developing.

Now open a web browser and enter the URL `http://localhost:8000/`, or `http://127.0.0.1:8000/` if the first one doesn't work. You should see something like Figure 18-1, a page that Django creates to let you know all is working properly so far. Keep the server running for now, but when you want to stop the server you can do so by pressing CTRL-C.




Figure 18-1: Everything is working so far.

**NOTE**

If you receive the error message *That port is already in use*, tell Django to use a different port by entering `python manage.py runserver 8001` and cycle through higher numbers until you find an open port.

**TRY IT YOURSELF**

**18-1. New Projects:** To get a better idea of what Django does, build a couple of empty projects and look at what it creates. Make a new folder with a simple name, like *InstaBook* or *FaceGram* (outside of your *learning\_log* directory), navigate to that folder in a terminal, and create a virtual environment. Install Django, and run the command `django-admin.py startproject instabook .` (make sure you include the dot at the end of the command).

Look at the files and folders this command creates, and compare them to Learning Log. Do this a few times until you're familiar with what Django creates when starting a new project. Then delete the project directories if you wish.

## Starting an App

A Django *project* is organized as a group of individual *apps* that work together to make the project work as a whole. For now, we'll create just one app to do most of the work for our project. We'll add another app to manage user accounts in Chapter 19.

You should still be running runserver in the terminal window you opened earlier. Open a new terminal window (or tab) and navigate to the directory that contains *manage.py*. Activate the virtual environment, and then run the startapp command:

---

```
learning_log$ source ll_env/bin/activate
(ll_env)learning_log$ python manage.py startapp learning_logs
❶ (ll_env)learning_log$ ls
db.sqlite3  learning_log  learning_logs  ll_env  manage.py
❷ (ll_env)learning_log$ ls learning_logs/
admin.py  __init__.py  migrations  models.py  tests.py  views.py
```

---

The command `startapp appname` tells Django to create the infrastructure needed to build an app. If you look in the project directory now, you'll see a new folder called *learning\_logs* ❶. Open that folder to see what Django has created ❷. The most important files are *models.py*, *admin.py*, and *views.py*. We'll use *models.py* to define the data we want to manage in our app. We'll get to *admin.py* and *views.py* a little later.

### Defining Models

Let's think about our data for a moment. Each user will need to create a number of topics in their learning log. Each entry they make will be tied to a topic, and these entries will be displayed as text. We'll also need to store the timestamp of each entry so we can show users when they made each entry.

Open the file *models.py*, and look at its existing content:

---

```
models.py  from django.db import models

# Create your models here.
```

---

A module called *models* is being imported for us, and we're being invited to create models of our own. A model tells Django how to work with the data that will be stored in the app. Code-wise, a model is just a class; it has attributes and methods, just like every class we've discussed. Here's the model for the topics users will store:

---

```
from django.db import models

class Topic(models.Model):
    """A topic the user is learning about"""
    ❶ text = models.CharField(max_length=200)
    ❷ date_added = models.DateTimeField(auto_now_add=True)
```

---

```
❸ def __str__(self):
    """Return a string representation of the model."""
    return self.text
```

---

We've created a class called `Topic`, which inherits from `Model`—a parent class included in Django that defines the basic functionality of a model. Only two attributes are in the `Topic` class: `text` and `date_added`.

The `text` attribute is a `CharField`—a piece of data that's made up of characters, or text ❶. You use `CharField` when you want to store a small amount of text, such as a name, a title, or a city. When we define a `CharField` attribute, we have to tell Django how much space it should reserve in the database. Here we give it a `max_length` of 200 characters, which should be enough to hold most topic names.

The `date_added` attribute is a `DateTimeField`—a piece of data that will record a date and time ❷. We pass the argument `auto_add_now=True`, which tells Django to automatically set this attribute to the current date and time whenever the user creates a new topic.

**NOTE** *To see the different kinds of fields you can use in a model, see the Django Model Field Reference at <https://docs.djangoproject.com/en/1.8/ref/models/fields/>. You won't need all the information right now, but it will be extremely useful when you're developing your own apps.*

We need to tell Django which attribute to use by default when it displays information about a topic. Django calls a `__str__()` method to display a simple representation of a model. Here we've written a `__str__()` method that returns the string stored in the `text` attribute ❸.

**NOTE** *If you're using Python 2.7, you should call the `__str__()` method `__unicode__()` instead. The body of the method is identical.*

## Activating Models

To use our models, we have to tell Django to include our app in the overall project. Open `settings.py` (in the `learning_log/learning_log` directory), and you'll see a section that tells Django which apps are installed in the project:

```
settings.py  --snip--
INSTALLED_APPS = (
    'django.contrib.admin',
    'django.contrib.auth',
    'django.contrib.contenttypes',
    'django.contrib.sessions',
    'django.contrib.messages',
    'django.contrib.staticfiles',
)
--snip--
```

---

This is just a tuple, telling Django which apps work together to make up the project. Add our app to this tuple by modifying `INSTALLED_APPS` so it looks like this:

---

```
--snip--
INSTALLED_APPS = (
    --snip--
    'django.contrib.staticfiles',

    # My apps
    'learning_logs',
)
--snip--
```

---

Grouping apps together in a project helps to keep track of them as the project grows to include more apps. Here we start a section called *My apps*, which includes only `learning_logs` for now.

Next, we need to tell Django to modify the database so it can store information related to the model `Topic`. From the terminal, run the following command:

---

```
(ll_env)learning_log$ python manage.py makemigrations learning_logs
Migrations for 'learning_logs':
  0001_initial.py:
    - Create model Topic
(ll_env)learning_log$
```

---

The command `makemigrations` tells Django to figure out how to modify the database so it can store the data associated with any new models we've defined. The output here shows that Django has created a migration file called *0001\_initial.py*. This migration will create a table for the model `Topic` in the database.

Now we'll apply this migration and have Django modify the database for us:

---

```
(ll_env)learning_log$ python manage.py migrate
--snip--
Running migrations:
  Rendering model states... DONE
❶ Applying learning_logs.0001_initial... OK
```

---

Most of the output from this command is identical to the output from the first time we issued the `migrate` command. The line we need to check appears at ❶, where Django confirms that everything worked OK when it applied the migration for `learning_logs`.

Whenever we want to modify the data that `Learning Log` manages, we'll follow these three steps: modify *models.py*, call `makemigrations` on `learning_logs`, and tell Django to migrate the project.

## The Django Admin Site

When you define models for an app, Django makes it easy for you to work with your models through the *admin site*. A site's administrators use the admin site, not a site's general users. In this section, we'll set up the admin site and use it to add some topics through the Topic model.

### Setting Up a Superuser

Django allows you to create a user who has all privileges available on the site, called a *superuser*. A *privilege* controls the actions a user can take. The most restrictive privilege settings allow a user to only read public information on the site. Registered users typically have the privilege of reading their own private data and some selected information available only to members. To effectively administer a web application, the site owner usually needs access to all information stored on the site. A good administrator is careful with their users' sensitive information, because users put a lot of trust into the apps they access.

To create a superuser in Django, enter the following command and respond to the prompts:

---

```
(ll_env)learning_log$ python manage.py createsuperuser
❶ Username (leave blank to use 'ehmatthes'): ll_admin
❷ Email address:
❸ Password:
Password (again):
Superuser created successfully.
(ll_env)learning_log$
```

---

When you issue the command `createsuperuser`, Django prompts you to enter a username for the superuser ❶. Here we're using `ll_admin`, but you can enter any username you want. You can enter an email address if you want or just leave this field blank ❷. You'll need to enter your password twice ❸.

#### NOTE

*Some sensitive information can be hidden from a site's administrators. For example, Django doesn't actually store the password you enter; instead, it stores a string derived from the password, called a hash. Each time you enter your password, Django hashes your entry and compares it to the stored hash. If the two hashes match, you're authenticated. By requiring hashes to match, if an attacker gains access to a site's database, they'll be able to read its stored hashes but not the passwords. When a site is set up properly, it's almost impossible to get the original passwords from the hashes.*

### Registering a Model with the Admin Site

Django includes some models in the admin site automatically, such as User and Group, but the models we create need to be registered manually.

When we started the `learning_logs` app, Django created a file called `admin.py` in the same directory as `models.py`:

```
admin.py from django.contrib import admin

# Register your models here.
```

To register `Topic` with the admin site, enter:

```
from django.contrib import admin

❶ from learning_logs.models import Topic

❷ admin.site.register(Topic)
```

This code imports the model we want to register, `Topic` ❶, and then uses `admin.site.register()` ❷ to tell Django to manage our model through the admin site.

Now use the superuser account to access the admin site. Go to `http://localhost:8000/admin/`, enter the username and password for the superuser you just created, and you should see a screen like the one in Figure 18-2. This page allows you to add new users and groups and change existing ones. We can also work with data related to the `Topic` model that we just defined.

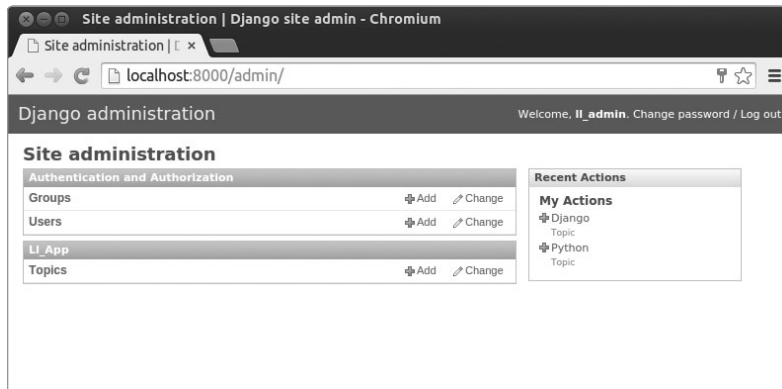


Figure 18-2: The admin site with `Topic` included

#### NOTE

If you see a message in your browser that the web page is not available, make sure you still have the Django server running in a terminal window. If you don't, activate a virtual environment and reissue the command `python manage.py runserver`.

## Adding Topics

Now that `Topic` has been registered with the admin site, let's add our first topic. Click **Topics** to go to the Topics page, which is mostly empty, because we have no topics to manage yet. Click **Add**, and you'll see a form for adding

a new topic. Enter **Chess** in the first box and click **Save**. You'll be sent back to the Topics admin page, and you'll see the topic you just created.

Let's create a second topic so we'll have more data to work with. Click **Add** again, and create a second topic, **Rock Climbing**. When you click **Save**, you'll be sent back to the main Topics page again, and you'll see both Chess and Rock Climbing listed.

## Defining the Entry Model

To record what we've been learning about chess and rock climbing, we need to define a model for the kinds of entries users can make in their learning logs. Each entry needs to be associated with a particular topic. This relationship is called a *many-to-one relationship*, meaning many entries can be associated with one topic.

Here's the code for the Entry model:

---

```
models.py  from django.db import models

class Topic(models.Model):
    --snip--

❶ class Entry(models.Model):
    """Something specific learned about a topic"""
    ❷ topic = models.ForeignKey(Topic)
    ❸ text = models.TextField()
    date_added = models.DateTimeField(auto_now_add=True)

    ❹ class Meta:
        verbose_name_plural = 'entries'

    def __str__(self):
        """Return a string representation of the model."""
    ❺ return self.text[:50] + "..."
```

---

The Entry class inherits from Django's base Model class, just as Topic did ❶. The first attribute, topic, is a ForeignKey instance ❷. A *foreign key* is a database term; it's a reference to another record in the database. This is the code that connects each entry to a specific topic. Each topic is assigned a key, or ID, when it's created. When Django needs to establish a connection between two pieces of data, it uses the key associated with each piece of information. We'll use these connections shortly to retrieve all the entries associated with a certain topic.

Next is an attribute called text, which is an instance of TextField ❸. This kind of field doesn't need a size limit, because we don't want to limit the size of individual entries. The date\_added attribute allows us to present entries in the order they were created and to place a timestamp next to each entry.

At ❹ we nest the Meta class inside our Entry class. Meta holds extra information for managing a model; here it allows us to set a special attribute telling Django to use *Entries* when it needs to refer to more than one entry.

(Without this, Django would refer to multiple entries as *Entrys*.) Finally, the `__str__()` method tells Django which information to show when it refers to individual entries. Because an entry can be a long body of text, we tell Django to show just the first 50 characters of text ❹. We also add an ellipsis to clarify that we're not always displaying the entire entry.

## Migrating the Entry Model

Because we've added a new model, we need to migrate the database again. This process will become quite familiar: you modify *models.py*, run the command `python manage.py makemigrations app_name`, and then run the command `python manage.py migrate`.

Migrate the database and check the output:

---

```
(ll_env)learning_log$ python manage.py makemigrations learning_logs
Migrations for 'learning_logs':
❶ 0002_entry.py:
    - Create model Entry
(ll_env)learning_log$ python manage.py migrate
Operations to perform:
  --snip--
❷ Applying learning_logs.0002_entry... OK
```

---

A new migration called *0002\_entry.py* is generated, which tells Django how to modify the database to store information related to the model Entry ❶. When we issue the migrate command, we see that Django applied this migration, and everything was okay ❷.

## Registering Entry with the Admin Site

We also need to register the Entry model. Here's what *admin.py* should look like now:

---

```
admin.py  from django.contrib import admin

         from learning_logs.models import Topic, Entry

         admin.site.register(Topic)
         admin.site.register(Entry)
```

---

Go back to <http://localhost/admin/>, and you should see *Entries* listed under *learning\_logs*. Click the **Add** link for Entries, or click **Entries**, and then choose **Add entry**. You should see a drop-down list to select the topic you're creating an entry for and a text box for adding an entry. Select **Chess** from the drop-down list, and add an entry. Here's the first entry I made:

The opening is the first part of the game, roughly the first ten moves or so. In the opening, it's a good idea to do three things—bring out your bishops and knights, try to control the center of the board, and castle your king.



Of course, these are just guidelines. It will be important to learn when to follow these guidelines and when to disregard these suggestions.

When you click **Save**, you'll be brought back to the main admin page for entries. Here you'll see the benefit of using `text[:50]` as the string representation for each entry; it's much easier to work with multiple entries in the admin interface if you see only the first part of an entry rather than the entire text of each entry.

Make a second entry for Chess and one entry for Rock Climbing so we have some initial data. Here's a second entry for Chess:

In the opening phase of the game, it's important to bring out your bishops and knights. These pieces are powerful and maneuverable enough to play a significant role in the beginning moves of a game.

And here's a first entry for Rock Climbing:

One of the most important concepts in climbing is to keep your weight on your feet as much as possible. There's a myth that climbers can hang all day on their arms. In reality, good climbers have practiced specific ways of keeping their weight over their feet whenever possible.

These three entries will give us something to work with as we continue to develop Learning Log.

## ***The Django Shell***

Now that we've entered some data, we can examine that data programmatically through an interactive terminal session. This interactive environment is called the Django *shell*, and it's a great environment for testing and troubleshooting your project. Here's an example of an interactive shell session:

---

```
(ll_env)learning_log$ python manage.py shell
❶ >>> from learning_logs.models import Topic
>>> Topic.objects.all()
[<Topic: Chess>, <Topic: Rock Climbing>]
```

---

The command `python manage.py shell` (run in an active virtual environment) launches a Python interpreter that you can use to explore the data stored in your project's database. Here we import the model `Topic` from the `learning_logs.models` module ❶. We then use the method `Topic.objects.all()` to get all of the instances of the model `Topic`; the list that's returned is called a *queryset*.

We can loop over a queryset just as we'd loop over a list. Here's how you can see the ID that's been assigned to each topic object:

---

```
>>> topics = Topic.objects.all()
>>> for topic in topics:
```

---

```
... print(topic.id, topic)
...
1 Chess
2 Rock Climbing
```

---

We store the queryset in `topics`, and then print each topic's `id` attribute and the string representation of each topic. We can see that Chess has an ID of 1, and Rock Climbing has an ID of 2.

If you know the ID of a particular object, you can get that object and examine any attribute the object has. Let's look at the `text` and `date_added` values for Chess:

---

```
>>> t = Topic.objects.get(id=1)
>>> t.text
'Chess'
>>> t.date_added
datetime.datetime(2015, 5, 28, 4, 39, 11, 989446, tzinfo=<UTC>)
```

---

We can also look at the entries related to a certain topic. Earlier we defined the `topic` attribute for the `Entry` model. This was a `ForeignKey`, a connection between each entry and a topic. Django can use this connection to get every entry related to a certain topic, like this:

---

```
❶ >>> t.entry_set.all()
[<Entry: The opening is the first part of the game, roughly...>, <Entry: In the opening phase of the game, it's important t...>]
```

---

To get data through a foreign key relationship, you use the lowercase name of the related model followed by an underscore and the word `set` ❶. For example, say you have the models `Pizza` and `Topping`, and `Topping` is related to `Pizza` through a foreign key. If your object is called `my_pizza`, representing a single pizza, you can get all of the pizza's toppings using the code `my_pizza.topping_set.all()`.

We'll use this kind of syntax when we begin to code the pages users can request. The shell is very useful for making sure your code retrieves the data you want it to. If your code works as you expect it to in the shell, you can expect it to work properly in the files you write within your project. If your code generates errors or doesn't retrieve the data you expect it to, it's much easier to troubleshoot your code in the simple shell environment than it is within the files that generate web pages. We won't refer to the shell much, but you should continue using it to practice working with Django's syntax for accessing the data stored in the project.

**NOTE**

*Each time you modify your models, you'll need to restart the shell to see the effects of those changes. To exit a shell session, enter `CTRL-D`; on Windows enter `CTRL-Z` and then press `ENTER`.*

### TRY IT YOURSELF

**18-2. Short Entries:** The `__str__()` method in the Entry model currently appends an ellipsis to every instance of Entry when Django shows it in the admin site or the shell. Add an if statement to the `__str__()` method that adds an ellipsis only if the entry is more than 50 characters long. Use the admin site to add an entry that's fewer than 50 characters in length, and check that it doesn't have an ellipsis when viewed.

**18-3. The Django API:** When you write code to access the data in your project, you're writing a *query*. Skim through the documentation for querying your data at <https://docs.djangoproject.com/en/1.8/topics/db/queries/>. Much of what you see will look new to you, but it will be quite useful as you start to work on your own projects.

**18-4. Pizzeria:** Start a new project called *pizzeria* with an app called *pizzas*. Define a model *Pizza* with a field called *name*, which will hold name values such as *Hawaiian* and *Meat Lovers*. Define a model called *Topping* with fields called *pizza* and *name*. The *pizza* field should be a foreign key to *Pizza*, and *name* should be able to hold values such as *pineapple*, *Canadian bacon*, and *sausage*.

Register both models with the admin site, and use the site to enter some pizza names and toppings. Use the shell to explore the data you entered.

## Making Pages: The Learning Log Home Page

Usually, making web pages with Django consists of three stages: defining URLs, writing views, and writing templates. First, you must define patterns for URLs. A URL pattern describes the way the URL is laid out and tells Django what to look for when matching a browser request with a site URL so it knows which page to return.

Each URL then maps to a particular *view*—the view function retrieves and processes the data needed for that page. The view function often calls a *template*, which builds a page that a browser can read. To see how this works, let's make the home page for Learning Log. We'll define the URL for the home page, write its view function, and create a simple template.

Because all we're doing is making sure Learning Log works as it's supposed to, we'll keep the page simple for now. A functioning web app is fun to style when it's complete; an app that looks good but doesn't work well is pointless. For now, the home page will display only a title and a brief description.

## Mapping a URL

Users request pages by entering URLs into a browser and clicking links, so we'll need to decide what URLs are needed in our project. The home page URL is first: it's the base URL people use to access the project. At the moment, the base URL, `http://localhost:8000/`, returns the default Django site that lets us know the project was set up correctly. We'll change this by mapping the base URL to Learning Log's home page.

In the main *learning\_log* project folder, open the file *urls.py*. Here's the code you'll see:

---

```
urls.py ❶ from django.conf.urls import include, url
        from django.contrib import admin

        ❷ urlpatterns = [
        ❸     url(r'^admin/', include(admin.site.urls)),
        ]
```

---

The first two lines import the functions and modules that manage URLs for the project and admin site ❶. The body of the file defines the `urlpatterns` variable ❷. In this *urls.py* file, which represents the project as a whole, the `urlpatterns` variable includes sets of URLs from the apps in the project. The code at ❸ includes the module `admin.site.urls`, which defines all the URLs that can be requested from the admin site.

We need to include the URLs for *learning\_logs*:

---

```
from django.conf.urls import include, url
from django.contrib import admin

urlpatterns = [
    url(r'^admin/', include(admin.site.urls)),
    ❶ url(r'', include('learning_logs.urls', namespace='learning_logs')),
]
```

---

We've added a line to include the module `learning_logs.urls` at ❶. This line includes a namespace argument, which allows us to distinguish *learning\_logs*'s URLs from other URLs that might appear in the project, which can be very helpful as your project starts to grow.

The default *urls.py* is in the *learning\_log* folder; now we need to make a second *urls.py* file in the *learning\_logs* folder:

---

```
urls.py ❶ """Defines URL patterns for learning_logs."""

        ❷ from django.conf.urls import url

        ❸ from . import views

        ❹ urlpatterns = [
            # Home page
        ❺ url(r'^$', views.index, name='index'),
        ]
```

---

To make it clear which `urls.py` we're working in, we add a docstring at the beginning of the file ❶. We then import the `url` function, which is needed when mapping URLs to views ❷. We also import the `views` module ❸; the dot tells Python to import `views` from the same directory as the current `urls.py` module. The variable `urlpatterns` in this module is a list of individual pages that can be requested from the `learning_logs` app ❹.

The actual URL pattern is a call to the `url()` function, which takes three arguments ❺. The first is a regular expression. Django will look for a regular expression in `urlpatterns` that matches the requested URL string. Therefore, a regular expression will define the pattern that Django can look for.

Let's look at the regular expression `r'^$',`. The `r` tells Python to interpret the following string as a raw string, and the quotes tell Python where the regular expression begins and ends. The caret (^) tells Python to find the beginning of the string, and the dollar sign tells Python to look for the end of the string. In its entirety, this expression tells Python to look for a URL with nothing between the beginning and end of the URL. Python ignores the base URL for the project (`http://localhost:8000/`), so an empty regular expression matches the base URL. Any other URL will not match this expression, and Django will return an error page if the URL requested doesn't match any existing URL patterns.

The second argument in `url()` at ❻ specifies which view function to call. When a requested URL matches the regular expression, Django will call `views.index` (we'll write this view function in the next section). The third argument provides the name `index` for this URL pattern so we can refer to it in other sections of the code. Whenever we want to provide a link to the home page, we'll use this name instead of writing out a URL.

#### NOTE

*Regular expressions, often called regexes, are used in almost every programming language. They're incredibly useful, but they take some practice to get used to. If you didn't follow all of this, don't worry; you'll see plenty of examples as you work through this project.*

## Writing a View

A view function takes in information from a request, prepares the data needed to generate a page, and then sends the data back to the browser, often by using a template that defines what the page will look like.

The file `views.py` in `learning_logs` was generated automatically when we ran the command `python manage.py startapp`. Here's what's in `views.py` right now:

---

```
views.py  from django.shortcuts import render

# Create your views here.
```

---

Currently, this file just imports the `render()` function, which renders the response based on the data provided by views. The following code is how the view for the home page should be written:

---

```
from django.shortcuts import render

def index(request):
    """The home page for Learning Log"""
    return render(request, 'learning_logs/index.html')
```

---

When a URL request matches the pattern we just defined, Django will look for a function called `index()` in the `views.py` file. Django then passes the request object to this view function. In this case, we don't need to process any data for the page, so the only code in the function is a call to `render()`. The `render()` function here uses two arguments—the original request object and a template it can use to build the page. Let's write this template.

## Writing a Template

A template sets up the structure for a web page. The template defines what the page should look like, and Django fills in the relevant data each time the page is requested. A template allows you to access any data provided by the view. Because our view for the home page provided no data, this template is fairly simple.

Inside the `learning_logs` folder, make a new folder called `templates`. Inside the `templates` folder, make another folder called `learning_logs`. This might seem a little redundant (we have a folder named `learning_logs` inside a folder named `templates` inside a folder named `learning_logs`), but it sets up a structure that Django can interpret unambiguously, even in the context of a large project containing many individual apps. Inside the inner `learning_logs` folder, make a new file called `index.html`. Write the following into that file:

---

```
index.html    <p>Learning Log</p>

               <p>Learning Log helps you keep track of your learning, for any topic you're
               learning about.</p>
```

---

This is a very simple file. If you're not familiar with HTML, the `<p></p>` tags signify paragraphs. The `<p>` tag opens a paragraph, and the `</p>` tag closes a paragraph. We have two paragraphs: the first acts as a title, and the second describes what users can do with Learning Log.

Now when we request the project's base URL, `http://localhost:8000/`, we'll see the page we just built instead of the default Django page. Django will take the requested URL, and that URL will match the pattern `r'^$',`; then Django will call the function `views.index()`, and this will render the page using the template contained in `index.html`. The resulting page is shown in Figure 18-3.

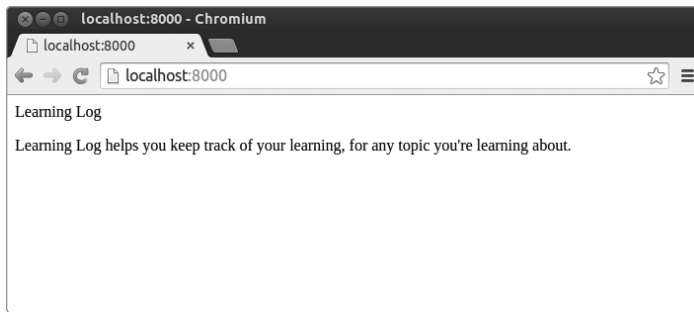


Figure 18-3: The home page for Learning Log

Although it may seem a complicated process for creating one page, this separation between URLs, views, and templates actually works well. It allows you to think about each aspect of a project separately, and in larger projects it allows individuals to focus on the areas in which they're strongest. For example, a database specialist can focus on the models, a programmer can focus on the view code, and a web designer can focus on the templates.

#### TRY IT YOURSELF

**18-5. Meal Planner:** Consider an app that helps people plan their meals throughout the week. Make a new folder called *meal\_planner*, and start a new Django project inside this folder. Then make a new app called *meal\_plans*. Make a simple home page for this project.

**18-6. Pizzeria Home Page:** Add a home page to the *Pizzeria* project you started in Exercise 18-4 (page 412).

## Building Additional Pages

Now that we've established a routine for building a page, we can start to build out the Learning Log project. We'll build two pages that display data: a page that lists all topics and a page that shows all the entries for a particular topic. For each of these pages, we'll specify a URL pattern, write a view function, and write a template. But before we do this, we'll create a base template that all templates in the project can inherit from.

### *Template Inheritance*

When building a website, you'll almost always require some elements to be repeated on each page. Rather than writing these elements directly into each page, you can write a base template containing the repeated elements

and then have each page inherit from the template. This approach lets you focus on developing the unique aspects of each page and makes it much easier to change the overall look and feel of the project.

## The Parent Template

We'll start by creating a template called *base.html* in the same directory as *index.html*. This file will contain elements common to all pages; every other template will inherit from *base.html*. The only element we want to repeat on each page right now is the title at the top. Because we'll include this template on every page, let's make the title a link to the home page:

---

```
base.html  <p>
           ❶  <a href="{% url 'learning_logs:index' %}">Learning Log</a>
           </p>

           ❷  {% block content %}{% endblock content %}
```

---

The first part of this file creates a paragraph containing the name of the project, which also acts as a link to the home page. To generate a link, we use a *template tag*, indicated by braces and percent signs {% %}. A template tag is a bit of code that generates information to be displayed on a page. In this example, the template tag {% url 'learning\_logs:index' %} generates a URL matching the URL pattern defined in *learning\_logs/urls.py* with the name 'index' ❶. In this example, *learning\_logs* is the *namespace* and *index* is a uniquely named URL pattern in that namespace.

In a simple HTML page, a link is surrounded by the *anchor* tag:

---

```
<a href="link_url">link text</a>
```

---

Having the template tag generate the URL for us makes it much easier to keep our links up to date. To change a URL in our project, we only need to change the URL pattern in *urls.py*, and Django will automatically insert the updated URL the next time the page is requested. Every page in our project will inherit from *base.html*, so from now on every page will have a link back to the home page.

At ❷ we insert a pair of block tags. This block, named *content*, is a placeholder; the child template will define the kind of information that goes in the content block.

A child template doesn't have to define every block from its parent, so you can reserve space in parent templates for as many blocks as you like, and the child template uses only as many as it requires.

### NOTE

*In Python code, we almost always indent four spaces. Template files tend to have more levels of nesting than Python files, so it's common to use only two spaces for each indentation level.*



## The Child Template

Now we need to rewrite *index.html* to inherit from *base.html*. Here's *index.html*:

---

```
index.html ❶ {% extends "learning_logs/base.html" %}

❷ {% block content %}
    <p>Learning Log helps you keep track of your learning, for any topic you're
    learning about.</p>
❸ {% endblock content %}
```

---

If you compare this to the original *index.html*, you can see that we've replaced the Learning Log title with the code for inheriting from a parent template ❶. A child template must have an `{% extends %}` tag on the first line to tell Django which parent template to inherit from. The file *base.html* is part of *learning\_logs*, so we include *learning\_logs* in the path to the parent template. This line pulls in everything contained in the *base.html* template and allows *index.html* to define what goes in the space reserved by the content block.

We define the content block at ❷ by inserting a `{% block %}` tag with the name *content*. Everything that we aren't inheriting from the parent template goes inside a content block. Here, that's the paragraph describing the Learning Log project. At ❸ we indicate that we're finished defining the content by using an `{% endblock content %}` tag.

You can start to see the benefit of template inheritance: in a child template we only need to include content that's unique to that page. This not only simplifies each template, but also makes it much easier to modify the site. To modify an element common to many pages, you only need to modify the element in the parent template. Your changes are then carried over to every page that inherits from that template. In a project that includes tens or hundreds of pages, this structure can make it much easier and faster to improve your site.

### NOTE

*In a large project, it's common to have one parent template called *base.html* for the entire site and parent templates for each major section of the site. All the section templates inherit from *base.html*, and each page in the site inherits from a section template. This way you can easily modify the look and feel of the site as a whole, any section in the site, or any individual page. This configuration provides a very efficient way to work, and it encourages you to steadily update your site over time.*

## The Topics Page

Now that we have an efficient approach to building pages, we can focus on our next two pages: the general topics page and the page to display entries for a single topic. The topics page will show all topics that users have created, and it's the first page that will involve working with data.

## The Topics URL Pattern

First, we define the URL for the topics page. It's common to choose a simple URL fragment that reflects the kind of information presented on the page. We'll use the word *topics*, so the URL `http://localhost:8000/topics/` will return this page. Here's how we modify `learning_logs/urls.py`:

---

```
urls.py    """Defines URL patterns for learning_logs."""
           --snip--
           urlpatterns = [
               # Home page
               url(r'^$', views.index, name='index'),

               # Show all topics.
               ❶ url(r'^topics/$', views.topics, name='topics'),
           ]
```

---

We've simply added `topics/` into the regular expression argument used for the home page URL ❶. When Django examines a requested URL, this pattern will match any URL that has the base URL followed by *topics*. You can include or omit a forward slash at the end, but there can't be anything else after the word *topics*, or the pattern won't match. Any request with a URL that matches this pattern will then be passed to the function `topics()` in `views.py`.

## The Topics View

The `topics()` function needs to get some data from the database and send it to the template. Here's what we need to add to `views.py`:

---

```
views.py   from django.shortcuts import render

           ❶ from .models import Topic

           def index(request):
               --snip--

           ❷ def topics(request):
               """Show all topics."""
               ❸ topics = Topic.objects.order_by('date_added')
               ❹ context = {'topics': topics}
               ❺ return render(request, 'learning_logs/topics.html', context)
```

---

We first import the model associated with the data we need ❶. The `topics()` function needs one parameter: the request object Django received from the server ❷. At ❸ we query the database by asking for the Topic objects, sorted by the `date_added` attribute. We store the resulting queryset in `topics`.

At ❹ we define a context that we'll send to the template. A *context* is a dictionary in which the keys are names we'll use in the template to access the data and the values are the data we need to send to the template. In this case, there's one key-value pair, which contains the set of topics we'll display on the page. When building a page that uses data, we pass the context variable to `render()` as well as the request object and the path to the template ❺.

## The Topics Template

The template for the topics page receives the context dictionary so the template can use the data that `topics()` provides. Make a file called *topics.html* in the same directory as *index.html*. Here's how we can display the topics in the template:

---

```
topics.html    {% extends "learning_logs/base.html" %}

               {% block content %}

                 <p>Topics</p>

❶    <ul>
❷      {% for topic in topics %}
❸        <li>{{ topic }}</li>
❹      {% empty %}
        <li>No topics have been added yet.</li>
❺      {% endfor %}
❻    </ul>

               {% endblock content %}
```

---

We start by using the `{% extends %}` tag to inherit from *base.html*, just as the index template does, and then open a content block. The body of this page contains a bulleted list of the topics that have been entered. In standard HTML, a bulleted list is called an *unordered list*, indicated by the tags `<ul></ul>`. We begin the bulleted list of topics at ❶.

At ❷ we have another template tag equivalent to a for loop, which loops through the list topics from the context dictionary. The code used in templates differs from Python in some important ways. Python uses indentation to indicate which lines of a for statement are part of a loop. In a template, every for loop needs an explicit `{% endfor %}` tag indicating where the end of the loop occurs. So in a template, you'll see loops written like this:

---

```
{% for item in list %}
    do something with each item
{% endfor %}
```

---

Inside the loop, we want to turn each topic into an item in the bulleted list. To print a variable in a template, wrap the variable name in double

braces. The code `{{ topic }}` at ❸ will be replaced by the value of `topic` on each pass through the loop. The braces won't appear on the page; they just indicate to Django that we're using a template variable. The HTML tag `<li></li>` indicates a list item. Anything between these tags, inside a pair of `<ul></ul>` tags, will appear as a bulleted item in the list.

At ❹ we use the `{% empty %}` template tag, which tells Django what to do if there are no items in the list. In this case, we print a message informing the user that no topics have been added yet. The last two lines close out the for loop ❺ and then close out the bulleted list ❻.

Now we need to modify the base template to include a link to the topics page:

---

```
base.html    <p>
❶    <a href="{% url 'learning_logs:index' %}">Learning Log</a> -
❷    <a href="{% url 'learning_logs:topics' %}">Topics</a>
    </p>

    {% block content %}{% endblock content %}
```

---

We add a dash after the link to the home page ❶, and then we add a link to the topics page, using the URL template tag again ❷. This line tells Django to generate a link matching the URL pattern with the name 'topics' in *learning\_logs/urls.py*.

Now when you refresh the home page in your browser, you'll see a *Topics* link. When you click the link, you'll see a page that looks similar to Figure 18-4.



Figure 18-4: The topics page

## Individual Topic Pages

Next, we need to create a page that can focus on a single topic, showing the topic name and all the entries for that topic. We'll again define a new URL pattern, write a view, and create a template. We'll also modify the topics page so each item in the bulleted list links to its corresponding topic page.

## The Topic URL Pattern

The URL pattern for the topic page is a little different than the other URL patterns we've seen so far because it will use the topic's `id` attribute to indicate which topic was requested. For example, if the user wants to see the detail page for the topic Chess, where the `id` is 1, the URL will be `http://localhost:8000/topics/1/`. Here's a pattern to match this URL, which goes in `learning_logs/urls.py`:

---

```
urls.py  --snip--
urlpatterns = [
    --snip--
    # Detail page for a single topic
    url(r'^topics/(?P<topic_id>\d+)/$', views.topic, name='topic'),
]
```

---

Let's examine the regular expression in this URL pattern, `r'^topics/(?P<topic_id>\d+)/$'`. The `r` tells Django to interpret the string as a raw string, and the expression is contained in quotes. The second part of the expression, `/(?P<topic_id>\d+)/`, matches an integer between two forward slashes and stores the integer value in an argument called `topic_id`. The parentheses surrounding this part of the expression captures the value stored in the URL; the `?P<topic_id>` part stores the matched value in `topic_id`; and the expression `\d+` matches any number of digits that appear between the forward slashes.

When Django finds a URL that matches this pattern, it calls the view function `topic()` with the value stored in `topic_id` as an argument. We'll use the value of `topic_id` to get the correct topic inside the function.

## The Topic View

The `topic()` function needs to get the topic and all associated entries from the database, as shown here:

---

```
views.py  --snip--
❶ def topic(request, topic_id):
    """Show a single topic and all its entries."""
    ❷ topic = Topic.objects.get(id=topic_id)
    ❸ entries = topic.entry_set.order_by('-date_added')
    ❹ context = {'topic': topic, 'entries': entries}
    ❺ return render(request, 'learning_logs/topic.html', context)
```

---

This is the first view function that requires a parameter other than the request object. The function accepts the value captured by the expression `(?P<topic_id>\d+)` and stores it in `topic_id` ❶. At ❷ we use `get()` to retrieve the topic, just as we did in the Django shell. At ❸ we get the entries associated with this topic, and we order them according to `date_added`: the minus sign in front of `date_added` sorts the results in reverse order, which will display the most recent entries first. We store the topic and entries in the context dictionary ❹ and send context to the template `topic.html` ❺.

**NOTE**

The code phrases at ❷ and ❸ are called queries, because they query the database for specific information. When you're writing queries like these in your own projects, it's very helpful to try them out in the Django shell first. You'll get much quicker feedback in the shell than you will by writing a view and template and then checking the results in a browser.

## The Topic Template

The template needs to display the name of the topic and the entries. We also need to inform the user if no entries have been made yet for this topic:

---

```
topic.html  {% extends 'learning_logs/base.html' %}

            {% block content %}

❶      <p>Topic: {{ topic }}</p>

            <p>Entries:</p>
❷      <ul>
❸      {% for entry in entries %}
            <li>
❹          <p>{{ entry.date_added|date:'M d, Y H:i' }}</p>
❺          <p>{{ entry.text|linebreaks }}</p>
            </li>
❻      {% empty %}
            <li>
                There are no entries for this topic yet.
            </li>
            {% endfor %}
        </ul>

    {% endblock content %}
```

---

We extend *base.html*, as we do for all pages in the project. Next, we show the topic that's currently being displayed ❶, which is stored in the template variable `{{ topic }}`. The variable `topic` is available because it's included in the context dictionary. We then start a bulleted list to show each of the entries ❷ and loop through them as we did the topics earlier ❸.

Each bullet will list two pieces of information: the timestamp and the full text of each entry. For the timestamp ❹, we display the value of the attribute `date_added`. In Django templates, a vertical line (`|`) represents a template *filter*—a function that modifies the value in a template variable. The filter `date: 'M d, Y H:i'` displays timestamps in the format *January 1, 2015 23:00*. The next line displays the full value of `text` rather than just the first 50 characters from `entry`. The filter `linebreaks` ❺ ensures that long text entries include line breaks in a format understood by browsers rather than showing a block of uninterrupted text. At ❻ we use the `{% empty %}` template tag to print a message informing the user that no entries have been made.

## Links from the Topics Page

Before we look at the topic page in a browser, we need to modify the topics template so each topic links to the appropriate page. Here's the change to *topics.html*:

*topics.html*

```
--snip--
{% for topic in topics %}
    <li>
        <a href="{% url 'learning_logs:topic' topic.id %}">{{ topic }}</a>
    </li>
{% empty %}
--snip--
```

We use the URL template tag to generate the proper link, based on the URL pattern in *learning\_logs* with the name 'topic'. This URL pattern requires a *topic\_id* argument, so we add the attribute *topic.id* to the URL template tag. Now each topic in the list of topics is a link to a topic page, such as *http://localhost:8000/topics/1/*.

If you refresh the topics page and click a topic, you should see a page that looks like Figure 18-5.

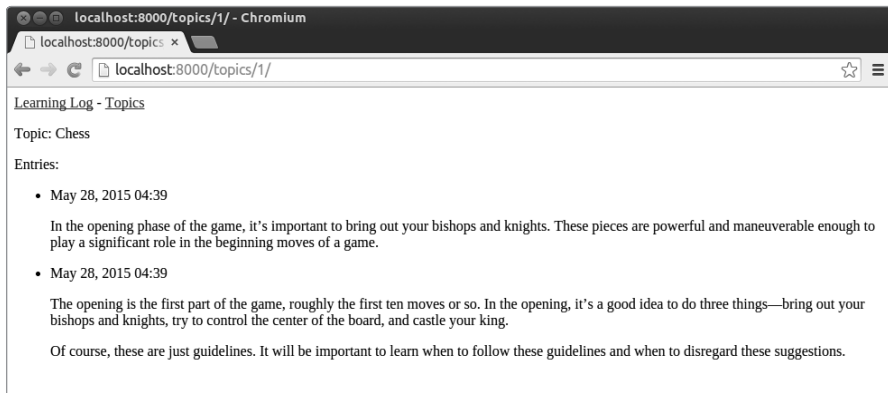


Figure 18-5: The detail page for a single topic, showing all entries for a topic

### TRY IT YOURSELF

**18-7. Template Documentation:** Skim the Django template documentation at <https://docs.djangoproject.com/en/1.8/ref/templates/>. You can refer back to it when you're working on your own projects.

**18-8. Pizzeria Pages:** Add a page to the *Pizzeria* project from Exercise 18-6 (page 416) that shows the names of available pizzas. Then link each pizza name to a page displaying the pizza's toppings. Make sure you use template inheritance to build your pages efficiently.

## Summary

In this chapter you started learning how to build web applications using the Django framework. You wrote a brief project spec, installed Django to a virtual environment, learned to set up a project, and checked that the project was set up correctly. You learned to set up an app and defined models to represent the data for your app. You learned about databases and how Django helps you migrate your database after you make a change to your models. You learned how to create a superuser for the admin site, and you used the admin site to enter some initial data.

You also explored the Django shell, which allows you to work with your project's data in a terminal session. You learned to define URLs, create view functions, and write templates to make pages for your site. Finally, you used template inheritance to simplify the structure of individual templates and to make it easier to modify the site as the project evolves.


In Chapter 19 we'll make intuitive, user-friendly pages that allow users to add new topics and entries and edit existing entries without going through the admin site. We'll also add a user registration system, allowing users to create an account and to make their own learning log. This is the heart of a web app—the ability to create something that any number of users can interact with.





# 19

## USER ACCOUNTS



At the heart of a web application is the ability for any user, anywhere in the world, to register an account with your app and start using it. In this chapter you'll build forms so users can add their own topics and entries, and edit existing entries. You'll also learn how Django guards against common attacks to form-based pages so you don't have to spend too much time thinking about securing your apps.

We'll then implement a user authentication system. You'll build a registration page for users to create accounts, and then restrict access to certain pages to logged-in users only. We'll then modify some of the view functions so users can only see their own data. You'll learn to keep your users' data safe and secure.

## Allowing Users to Enter Data

Before we build an authentication system for creating accounts, we'll first add some pages that allow users to enter their own data. We'll give users the ability to add a new topic, add a new entry, and edit their previous entries.

Currently, only a superuser can enter data through the admin site. We don't want users to interact with the admin site, so we'll use Django's form-building tools to build pages that allow users to enter data.

### Adding New Topics

Let's start by giving users the ability to add a new topic. Adding a form-based page works in much the same way as the pages we've already built: we define a URL, write a view function, and write a template. The one major difference is the addition of a new module called *forms.py*, which will contain the forms.

#### The Topic ModelForm

Any page that lets a user enter and submit information on a web page is a *form*, even if it doesn't look like one. When users enter information, we need to *validate* that the information provided is the right kind of data and not anything malicious, such as code to interrupt our server. We then need to process and save valid information to the appropriate place in the database. Django automates much of this work.

The simplest way to build a form in Django is to use a *ModelForm*, which uses the information from the models we defined in Chapter 18 to automatically build a form. Write your first form in the file *forms.py*, which you should create in the same directory as *models.py*:

---

```
forms.py  from django import forms

          from .models import Topic

❶ class TopicForm(forms.ModelForm):
          class Meta:
❷             model = Topic
❸             fields = ['text']
❹             labels = {'text': ''}
```

---

We first import the *forms* module and the model we'll work with, *Topic*. At ❶ we define a class called *TopicForm*, which inherits from *forms.ModelForm*.

The simplest version of a *ModelForm* consists of a nested *Meta* class telling Django which model to base the form on and which fields to include in the form. At ❷ we build a form from the *Topic* model and include only the text field ❸. The code at ❹ tells Django not to generate a label for the text field.

## The new\_topic URL

The URL for a new page should be short and descriptive, so when the user wants to add a new topic, we'll send them to *http://localhost:8000/new\_topic/*. Here's the URL pattern for the new\_topic page, which we add to *learning\_logs/urls.py*:

---

```
urls.py  --snip--
urlpatterns = [
    --snip--
    # Page for adding a new topic
    url(r'^new_topic/$', views.new_topic, name='new_topic'),
]
```

---

This URL pattern will send requests to the view function *new\_topic()*, which we'll write next.

## The new\_topic() View Function

The *new\_topic()* function needs to handle two different situations: initial requests for the new\_topic page (in which case it should show a blank form) and the processing of any data submitted in the form. It then needs to redirect the user back to the topics page:

---

```
views.py  from django.shortcuts import render
          from django.http import HttpResponseRedirect
          from django.core.urlresolvers import reverse

          from .models import Topic
          from .forms import TopicForm

          --snip--
          def new_topic(request):
              """Add a new topic."""
              ❶ if request.method != 'POST':
                  # No data submitted; create a blank form.
                  ❷ form = TopicForm()
              else:
                  # POST data submitted; process data.
                  ❸ form = TopicForm(request.POST)
                  ❹ if form.is_valid():
                      ❺ form.save()
                      ❻ return HttpResponseRedirect(reverse('learning_logs:topics'))

              ❼ context = {'form': form}
              return render(request, 'learning_logs/new_topic.html', context)
```

---

We import the class *HttpResponseRedirect*, which we'll use to redirect the reader back to the topics page after they submit their topic. The *reverse()* function determines the URL from a named URL pattern, meaning that Django will generate the URL when the page is requested. We also import the form we just wrote, *TopicForm*.

## GET and POST Requests

The two main types of request you'll use when building web apps are GET requests and POST requests. You use *GET* requests for pages that only read data from the server. You usually use *POST* requests when the user needs to submit information through a form. We'll be specifying the POST method for processing all of our forms. (A few other kinds of requests exist, but we won't be using them in this project.)

The function `new_topic()` takes in the request object as a parameter. When the user initially requests this page, their browser will send a GET request. When the user has filled out and submitted the form, their browser will submit a POST request. Depending on the request, we'll know whether the user is requesting a blank form (a GET request) or asking us to process a completed form (a POST request).

The test at ❶ determines whether the request method is GET or POST. If the request method is not POST, the request is probably GET, so we need to return a blank form (if it's another kind of request, it's still safe to return a blank form). We make an instance of `TopicForm` ❷, store it in the variable `form`, and send the form to the template in the context dictionary ❸. Because we included no arguments when instantiating `TopicForm`, Django creates a blank form that the user can fill out.

If the request method is POST, the `else` block runs and processes the data submitted in the form. We make an instance of `TopicForm` ❹ and pass it the data entered by the user, stored in `request.POST`. The `form` object that's returned contains the information submitted by the user.

We can't save the submitted information in the database until we've checked that it's valid ❺. The `is_valid()` function checks that all required fields have been filled in (all fields in a form are required by default) and that the data entered matches the field types expected—for example, that the length of text is less than 200 characters, as we specified in `models.py` in Chapter 18. This automatic validation saves us a lot of work. If everything is valid, we can call `save()` ❻, which writes the data from the form to the database. Once we've saved the data, we can leave this page. We use `reverse()` to get the URL for the topics page and pass the URL to `HttpResponseRedirect()` ❼, which redirects the user's browser to the topics page. On the topics page, the user should see the topic they just entered in the list of topics.

## The `new_topic` Template

Now we make a new template called `new_topic.html` to display the form we just created:

---

```
new_topic.html    {% extends "learning_logs/base.html" %}

                  {% block content %}
                    <p>Add a new topic:</p>
```

```

❶ <form action="{% url 'learning_logs:new_topic' %}" method='post'>
❷   {% csrf_token %}
❸   {{ form.as_p }}
❹   <button name="submit">add topic</button>
</form>

{% endblock content %}

```

---

This template extends *base.html*, so it has the same base structure as the rest of the pages in Learning Log. At ❶ we define an HTML form. The action argument tells the server where to send the data submitted in the form; in this case, we send it back to the view function *new\_topic()*. The method argument tells the browser to submit the data as a POST request.

Django uses the template tag `{% csrf_token %}` ❷ to prevent attackers from using the form to gain unauthorized access to the server (this kind of attack is called a *cross-site request forgery*). At ❸ we display the form; here you see how simple Django can make tasks such as displaying a form. We only need to include the template variable `{{ form.as_p }}` for Django to create all the fields necessary to display the form automatically. The `as_p` modifier tells Django to render all the form elements in paragraph format, which is a simple way to display the form neatly.

Django doesn't create a submit button for forms, so we define one at ❹.

### Linking to the new\_topic Page

Next, we include a link to the *new\_topic* page on the *topics* page:

```

topics.html  {% extends "learning_logs/base.html" %}

{% block content %}

    <p>Topics</p>

    <ul>
        --snip--
    </ul>

    <a href="{% url 'learning_logs:new_topic' %}">Add a new topic:</a>

{% endblock content %}

```

---

Place the link after the list of existing topics. Figure 19-1 shows the resulting form. Go ahead and use the form to add a few new topics of your own.




Figure 19-1: The page for adding a new topic

## Adding New Entries

Now that the user can add a new topic, they'll want to add new entries too. We'll again define a URL, write a view function and a template, and link to the page. But first we'll add another class to *forms.py*.

### The Entry ModelForm

We need to create a form associated with the Entry model, but this time with a little more customization than TopicForm:

---

```
forms.py  from django import forms

          from .models import Topic, Entry

          class TopicForm(forms.ModelForm):
              --snip--

          class EntryForm(forms.ModelForm):
              class Meta:
                  model = Entry
                  fields = ['text']
                  labels = {'text': ''}
                  widgets = {'text': forms.Textarea(attrs={'cols': 80})}
```

---

We first update the import statement to include Entry as well as Topic. The new class EntryForm inherits from `forms.ModelForm` and has a nested Meta class listing the model it's based on and the field to include in the form. We again give the field 'text' a blank label ❶.

At ❷ we include the widgets attribute. A *widget* is an HTML form element, such as a single-line text box, multi-line text area, or drop-down list. By including the widgets attribute you can override Django's default widget choices. By telling Django to use a `forms.Textarea` element, we're customizing the input widget for the field 'text' so the text area will be 80 columns wide instead of the default 40. This will give users enough room to write a meaningful entry.

## The new\_entry URL

We need to include a `topic_id` argument in the URL for adding a new entry, because the entry must be associated with a particular topic. Here's the URL, which we add to `learning_logs/urls.py`:

---

```
urls.py  --snip--
urlpatterns = [
    --snip--
    # Page for adding a new entry
    url(r'^new_entry/(?P<topic_id>\d+)/$', views.new_entry, name='new_entry'),
]
```

---

This URL pattern matches any URL with the form `http://localhost:8000/new_entry/id/`, where *id* is a number matching the topic ID. The code `(?P<topic_id>\d+)` captures a numerical value and stores it in the variable `topic_id`. When a URL matching this pattern is requested, Django sends the request and the ID of the topic to the `new_entry()` view function.

## The new\_entry() View Function

The view function for `new_entry` is much like the function for adding a new topic:

---

```
views.py  from django.shortcuts import render
          --snip--

          from .models import Topic
          from .forms import TopicForm, EntryForm

          --snip--
          def new_entry(request, topic_id):
              """Add a new entry for a particular topic."""
              ❶ topic = Topic.objects.get(id=topic_id)

              ❷ if request.method != 'POST':
                  # No data submitted; create a blank form.
                  ❸ form = EntryForm()
              else:
                  # POST data submitted; process data.
                  ❹ form = EntryForm(data=request.POST)
                  if form.is_valid():
                      ❺ new_entry = form.save(commit=False)
                      ❻ new_entry.topic = topic
                      new_entry.save()
                      ❼ return HttpResponseRedirect(reverse('learning_logs:topic',
                                                                  args=[topic_id]))

              context = {'topic': topic, 'form': form}
              return render(request, 'learning_logs/new_entry.html', context)
```

---



We update the `import` statement to include the `EntryForm` we just made. The definition of `new_entry()` has a `topic_id` parameter to store the value it receives from the URL. We'll need the topic to render the page and process the form's data, so we use `topic_id` to get the correct topic object at ❶.

At ❷ we check if the request method is POST or GET. The `if` block executes if it's a GET request, and we create a blank instance of `EntryForm` ❸. If the request method is POST, we process the data by making an instance of `EntryForm`, populated with the POST data from the request object ❹. We then check if the form is valid. If it is, we need to set the entry object's `topic` attribute before saving it to the database.

When we call `save()`, we include the argument `commit=False` ❺ to tell Django to create a new entry object and store it in `new_entry` without saving it to the database yet. We set `new_entry`'s `topic` attribute to the topic we pulled from the database at the beginning of the function ❻, and then we call `save()` with no arguments. This saves the entry to the database with the correct associated topic.

At ❼ we redirect the user to the topic page. The `reverse()` call requires two arguments—the name of the URL pattern we want to generate a URL for and an `args` list containing any arguments that need to be included in the URL. The `args` list has one item in it, `topic_id`. The `HttpResponseRedirect()` call then redirects the user to the topic page they made an entry for, and they should see their new entry in the list of entries.

### The `new_entry` Template

As you can see in the following code, the template for `new_entry` is similar to the template for `new_topic`:

---

```
new_entry.html    {% extends "learning_logs/base.html" %}

                 {% block content %}

❶    <p><a href="{% url 'learning_logs:topic' topic.id %}">{{ topic }}</a></p>

                 <p>Add a new entry:</p>
❷    <form action="{% url 'learning_logs:new_entry' topic.id %}" method='post'>
                 {% csrf_token %}
                 {{ form.as_p }}
                 <button name='submit'>add entry</button>
                 </form>

                 {% endblock content %}
```

---

We show the topic at the top of the page ❶, so the user can see which topic they're adding an entry to. This also acts as a link back to the main page for that topic.

The form's `action` argument includes the `topic_id` value in the URL, so the view function can associate the new entry with the correct topic ❷. Other than that, this template looks just like `new_topic.html`.

## Linking to the new\_entry Page

Next, we need to include a link to the `new_entry` page from each topic page:

*topic.html*

```
{% extends "learning_logs/base.html" %}

{% block content %}

    <p>Topic: {{ topic }}</p>

    <p>Entries:</p>
    <p>
        <a href="{% url 'learning_logs:new_entry' topic.id %}">add new entry</a>
    </p>
    <ul>
        --snip--
    </ul>

{% endblock content %}
```

We add the link just before showing the entries, because adding a new entry will be the most common action on this page. Figure 19-2 shows the `new_entry` page. Now users can add new topics and as many entries as they want for each topic. Try out the `new_entry` page by adding a few entries to some of the topics you've created.

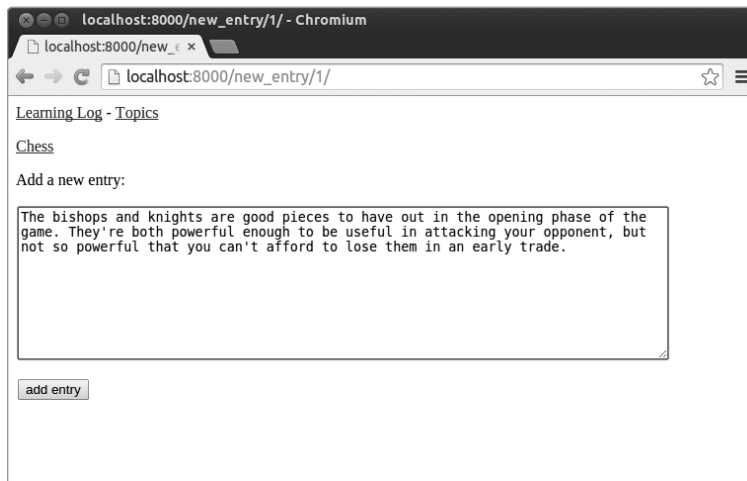


Figure 19-2: The `new_entry` page

## Editing Entries

Now we'll make a page to allow users to edit the entries they've already added.

## The edit\_entry URL

The URL for the page needs to pass the ID of the entry to be edited. Here's *learning\_logs/urls.py*:

---

```
urls.py  --snip--
urlpatterns = [
    --snip--
    # Page for editing an entry
    url(r'^edit_entry/(?P<entry_id>\d+)/$', views.edit_entry,
        name='edit_entry'),
]
```

---

The ID passed in the URL (for example, *http://localhost:8000/edit\_entry/1/*) is stored in the parameter *entry\_id*. The URL pattern sends requests that match this format to the view function *edit\_entry()*.

## The edit\_entry() View Function

When the *edit\_entry* page receives a GET request, *edit\_entry()* will return a form for editing the entry. When the page receives a POST request with revised entry text, it will save the modified text into the database:

---

```
views.py  from django.shortcuts import render
          --snip--

          from .models import Topic, Entry
          from .forms import TopicForm, EntryForm
          --snip--

          def edit_entry(request, entry_id):
              """Edit an existing entry."""
              ❶ entry = Entry.objects.get(id=entry_id)
                topic = entry.topic

                if request.method != 'POST':
                    # Initial request; pre-fill form with the current entry.
                    ❷ form = EntryForm(instance=entry)
                else:
                    # POST data submitted; process data.
                    ❸ form = EntryForm(instance=entry, data=request.POST)
                    if form.is_valid():
                        ❹ form.save()
                        ❺ return HttpResponseRedirect(reverse('learning_logs:topic',
                                                                    args=[topic.id]))

              context = {'entry': entry, 'topic': topic, 'form': form}
              return render(request, 'learning_logs/edit_entry.html', context)
```

---

We first need to import the *Entry* model. At ❶ we get the entry object that the user wants to edit and the topic associated with this entry. In the *if* block, which runs for a GET request, we make an instance of *EntryForm* with

the argument `instance=entry` ❷. This argument tells Django to create the form prefilled with information from the existing entry object. The user will see their existing data and be able to edit that data.

When processing a POST request, we pass the `instance=entry` argument and the `data=request.POST` argument ❸ to tell Django to create a form instance based on the information associated with the existing entry object, updated with any relevant data from `request.POST`. We then check if the form is valid; if it is, we call `save()` with no arguments ❹. We then redirect to the topic page ❺, where the user should see the updated version of the entry they edited.

### The `edit_entry` Template

Here's `edit_entry.html`, which is similar to `new_entry.html`:

---

```
edit_entry.html  {% extends "learning_logs/base.html" %}

                {% block content %}

                    <p><a href="{% url 'learning_logs:topic' topic.id %}">{{ topic }}</a></p>

                    <p>Edit entry:</p>

❶  <form action="{% url 'learning_logs:edit_entry' entry.id %}" method='post'>
    {% csrf_token %}
    {{ form.as_p }}
❷  <button name="submit">save changes</button>
    </form>

                {% endblock content %}
```

---

At ❶ the action argument sends the form back to the `edit_entry()` function for processing. We include the entry ID as an argument in the `{% url %}` tag, so the view function can modify the correct entry object. We label the submit button as *save changes* to remind the user they're saving edits, not creating a new entry ❷.

### Linking to the `edit_entry` Page

Now we need to include a link to the `edit_entry` page for each entry on the topic page:

---

```
topic.html      --snip--
                {% for entry in entries %}
                    <li>
                        <p>{{ entry.date_added|date:'M d, Y H:i' }}</p>
                        <p>{{ entry.text|linebreaks }}</p>
                        <p>
                            <a href="{% url 'learning_logs:edit_entry' entry.id %}">edit entry</a>
                        </p>
                    </li>
                {% endfor %}
                --snip--
```

---

We include the edit link after each entry's date and text has been displayed. We use the `{% url %}` template tag to determine the URL for the named URL pattern `edit_entry`, along with the ID attribute of the current entry in the loop (`entry.id`). The link text "edit entry" appears after each entry on the page. Figure 19-3 shows what the topic page looks like with these links.

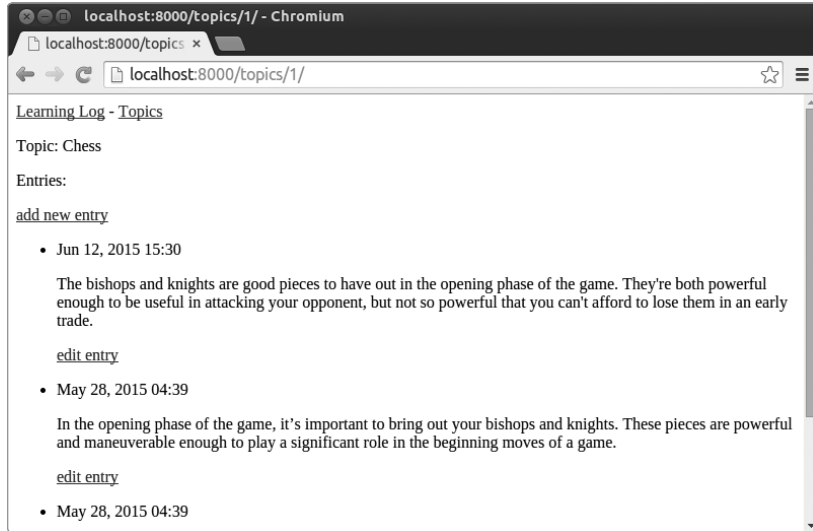


Figure 19-3: Each entry now has a link for editing that entry.

Learning Log now has most of the functionality it needs. Users can add topics and entries, and read through any set of entries they want. In the next section, we'll implement a user registration system so anyone can make an account with Learning Log and create their own set of topics and entries.

### TRY IT YOURSELF

**19-1. Blog:** Start a new Django project called *Blog*. Create an app called *blogs* in the project, with a model called *BlogPost*. The model should have fields like *title*, *text*, and *date\_added*. Create a superuser for the project, and use the admin site to make a couple of short posts. Make a home page that shows all posts in chronological order.

Create a form for making new posts and another for editing existing posts. Fill in your forms to make sure they work.

## Setting Up User Accounts

In this section we'll set up a user registration and authorization system to allow people to register an account and log in and out. We'll create a new app to contain all the functionality related to working with users. We'll also modify the Topic model slightly so every topic belongs to a certain user.

### The users App

We'll start by creating a new app called *users*, using the `startapp` command:

---

```
(ll_env)learning_log$ python manage.py startapp users
(ll_env)learning_log$ ls
❶ db.sqlite3 learning_log learning_logs ll_env manage.py users
(ll_env)learning_log$ ls users
❷ admin.py __init__.py migrations models.py tests.py views.py
```

---

This command makes a new directory called *users* ❶ with a structure identical to the *learning\_logs* app ❷.

### Adding users to settings.py

We need to add our new app to `INSTALLED_APPS` in *settings.py*, like so:

---

```
settings.py  --snip--
INSTALLED_APPS = (
    --snip--
    # My apps
    'learning_logs',
    'users',
)
--snip--
```

---

Now Django will include the *users* app in the overall project.

### Including the URLs from users

Next, we need to modify the root *urls.py* so it includes the URLs we'll write for the *users* app:

---

```
urls.py  from django.conf.urls import include, url
        from django.contrib import admin

        urlpatterns = [
            url(r'^admin/', include(admin.site.urls)),
            url(r'^users/', include('users.urls', namespace='users')),
            url(r'', include('learning_logs.urls', namespace='learning_logs')),
        ]
```

---

We add a line to include the file *urls.py* from *users*. This line will match any URL that starts with the word *users*, such as `http://localhost:8000/users/`

*login/*. We also create the namespace 'users' so we'll be able to distinguish URLs that belong to the *learning\_logs* app from URLs that belong to the *users* app.

## The Login Page

We'll first implement a login page. We'll use the default login view Django provides, so the URL pattern looks a little different. Make a new *urls.py* file in the directory *learning\_log/users/*, and add the following to it:

---

```
urls.py    """Defines URL patterns for users"""

    from django.conf.urls import url
    ❶ from django.contrib.auth.views import login

    from . import views

    urlpatterns = [
        # Login page
    ❷ url(r'^login/$', login, {'template_name': 'users/login.html'},
        name='login'),
    ]
```

---

We first import the default login view ❶. The login page's pattern matches the URL *http://localhost:8000/users/login/* ❷. When Django reads this URL, the word *users* tells Django to look in *users/urls.py*, and *login* tells it to send requests to Django's default login view (notice the view argument is *login*, not *views.login*). Because we're not writing our own view function, we pass a dictionary telling Django where to find the template we're about to write. This template will be part of the *users* app, not the *learning\_logs* app.

## The login Template

When the user requests the login page, Django will use its default login view, but we still need to provide a template for the page. Inside the *learning\_log/users/* directory, make a directory called *templates*; inside that, make another directory called *users*. Here's the *login.html* template, which you should save in *learning\_log/users/templates/users/*:

---

```
login.html    {% extends "learning_logs/base.html" %}

    {% block content %}

    ❶    {% if form.errors %}
        <p>Your username and password didn't match. Please try again.</p>
        {% endif %}

    ❷    <form method="post" action="{% url 'users:login' %}">
        {% csrf_token %}
    ❸    {{ form.as_p }}

    ❹    <button name="submit">log in</button>
```

```

❶ <input type="hidden" name="next" value="{% url 'learning_logs:index' %}" />
</form>

{% endblock content %}

```

---

This template extends *base.html* to ensure that the login page will have the same look and feel as the rest of the site. Note that a template in one app can extend a template from another app.

If the form's `errors` attribute is set, we display an error message ❶, reporting that the username and password combination don't match anything stored in the database.

We want the login view to process the form, so we set the action argument as the URL of the login page ❷. The login view sends a form to the template, and it's up to us to display the form ❸ and add a submit button ❹. At ❺ we include a hidden form element, 'next'; the value argument tells Django where to redirect the user after they've logged in successfully. In this case, we send the user back to the home page.

### Linking to the Login Page

Let's add the login link to *base.html* so it appears on every page. We don't want the link to display when the user is already logged in, so we nest it inside an `{% if %}` tag:

```

base.html <p>
          <a href="{% url 'learning_logs:index' %}">Learning Log</a> -
          <a href="{% url 'learning_logs:topics' %}">Topics</a> -
❶   {% if user.is_authenticated %}
❷   Hello, {{ user.username }}.
          {% else %}
❸   <a href="{% url 'users:login' %}">log in</a>
          {% endif %}
        </p>

{% block content %}{% endblock content %}

```

---

In Django's authentication system, every template has a user variable available, which always has an `is_authenticated` attribute set: the attribute is `True` if the user is logged in and `False` if they aren't. This allows you to display one message to authenticated users and another to unauthenticated users.

Here we display a greeting to users currently logged in ❶. Authenticated users have an additional `username` attribute set, which we use to personalize the greeting and remind the user they're logged in ❷. At ❸ we display a link to the login page for users who haven't been authenticated.

### Using the Login Page

We've already set up a user account, so let's log in to see if the page works. Go to `http://localhost:8000/admin/`. If you're still logged in as an admin, look for a logout link in the header and click it.



When you're logged out, go to `http://localhost:8000/users/login/`. You should see a login page similar to the one shown in Figure 19-4. Enter the username and password you set up earlier, and you should be brought back to the index page. The header on the home page should display a greeting personalized with your username.




Figure 19-4: The login page

## Logging Out

Now we need to provide a way for users to log out. We won't build a page for logging out; users will just click a link and be sent back to the home page. We'll define a URL pattern for the logout link, write a view function, and provide a logout link in *base.html*.

### The logout URL

The following code defines the URL pattern for logging out, matching the URL `http://localhost:8000/users/logout/`. Here's *users/urls.py*:

---

```
urls.py  --snip--
urlpatterns = [
    # Login page
    --snip--
    # Logout page
    url(r'^logout/$', views.logout_view, name='logout'),
]
```

---

The URL pattern sends the request to the `logout_view()` function, which is named as such to distinguish it from the `logout()` function we'll call from within the view. (Make sure you're modifying *users/urls.py*, not *learning\_log/urls.py*.)

### The `logout_view()` View Function

The `logout_view()` function is straightforward: we just import Django's `logout()` function, call it, and then redirect back to the home page. Open *users/views.py*, and enter the following code.

---

```
views.py  from django.http import HttpResponseRedirect
          from django.core.urlresolvers import reverse
          ❶ from django.contrib.auth import logout

          def logout_view(request):
              """Log the user out."""
              ❷ logout(request)
              ❸ return HttpResponseRedirect(reverse('learning_logs:index'))
```

---

We import the `logout()` function from `django.contrib.auth` ❶. In the function, we call `logout()` ❷, which requires the request object as an argument. We then redirect to the home page ❸.

### Linking to the logout View

Now we need a logout link. We'll include it as part of *base.html* so it's available on every page and include it in the `{% if user.is_authenticated %}` portion so only users who are already logged in can see it:

---

```
base.html  --snip--
           {% if user.is_authenticated %}
           Hello, {{ user.username }}.
           <a href="{% url 'users:logout' %}">log out</a>
           {% else %}
           <a href="{% url 'users:login' %}">log in</a>
           {% endif %}
           --snip--
```

---

Figure 19-5 shows the current home page as it appears to a logged-in user. The styling is minimal because we're focusing on building a site that works properly. When the required set of features works, we'll style the site to look more professional.

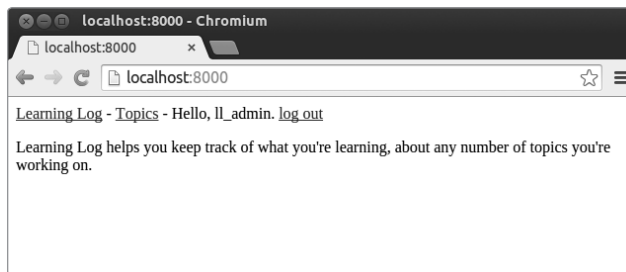


Figure 19-5: The home page with a personalized greeting and a logout link

### The Registration Page

Next, we'll build a page to allow new users to register. We'll use Django's default `UserCreationForm` but write our own view function and template.

## The register URL

The following code provides the URL pattern for the registration page, again in *users/urls.py*:

---

```
urls.py  --snip--
urlpatterns = [
    # Login page
    --snip--
    # Registration page
    url(r'^register/$', views.register, name='register'),
]
```

---

This pattern matches the URL *http://localhost:8000/users/register/* and sends requests to the *register()* function we're about to write.

## The register() View Function

The *register()* view function needs to display a blank registration form when the registration page is first requested and then process completed registration forms when they're submitted. When a registration is successful, the function also needs to log in the new user. Add the following code to *users/views.py*:

---

```
views.py  from django.shortcuts import render
          from django.http import HttpResponseRedirect
          from django.core.urlresolvers import reverse
          from django.contrib.auth import login, logout, authenticate
          from django.contrib.auth.forms import UserCreationForm

          def logout_view(request):
          --snip--

          def register(request):
              """Register a new user."""
              if request.method != 'POST':
                  # Display blank registration form.
                  ❶ form = UserCreationForm()
              else:
                  # Process completed form.
                  ❷ form = UserCreationForm(data=request.POST)

                  ❸ if form.is_valid():
                      ❹ new_user = form.save()
                      # Log the user in and then redirect to home page.
                      ❺ authenticated_user = authenticate(username=new_user.username,
                                                              password=request.POST['password1'])
                      ❻ login(request, authenticated_user)
                      ❼ return HttpResponseRedirect(reverse('learning_logs:index'))

              context = {'form': form}
              return render(request, 'users/register.html', context)
```

---

We first import the `render()` function. We then import the `login()` and `authenticate()` functions to log in the user if their registration information is correct. We also import the default `UserCreationForm`. In the `register()` function, we check whether or not we're responding to a POST request. If we're not, we make an instance of `UserCreationForm` with no initial data ❶.

If we're responding to a POST request, we make an instance of `UserCreationForm` based on the submitted data ❷. We check that the data is valid ❸—in this case, that the username has the appropriate characters, the passwords match, and the user isn't trying to do anything malicious in their submission.

If the submitted data is valid, we call the form's `save()` method to save the username and the hash of the password to the database ❹. The `save()` method returns the newly created user object, which we store in `new_user`.

When the user's information is saved, we log them in, which is a two-step process: we call `authenticate()` with the arguments `new_user.username` and their password ❺. When they register, the user is asked to enter two matching passwords, and because the form is valid, we know the passwords match so we can use either one. Here we get the value associated with the 'password1' key in the form's POST data. If the username and password are correct, the method returns an authenticated user object, which we store in `authenticated_user`. We then call the `login()` function with the request and `authenticated_user` objects ❻, which creates a valid session for the new user. Finally, we redirect the user to the home page ❼ where a personalized greeting in the header tells them their registration was successful.

## The register Template

The template for the registration page is similar to the login page. Be sure to save it in the same directory as *login.html*:

---

```
register.html  {% extends "learning_logs/base.html" %}

{% block content %}

    <form method="post" action="{% url 'users:register' %}">
        {% csrf_token %}
        {{ form.as_p }}

        <button name="submit">register</button>
        <input type="hidden" name="next" value="{% url 'learning_logs:index' %}" />
    </form>

{% endblock content %}
```

---

We use the `as_p` method again so Django will display all the fields in the form appropriately, including any error messages if the form is not filled out correctly.

## Linking to the Registration Page

Next, we'll add the code to show the registration page link to any user who is not currently logged in:

*base.html*

```
--snip--
{% if user.is_authenticated %}
    Hello, {{ user.username }}.
    <a href="{% url 'users:logout' %}">log out</a>
{% else %}
    <a href="{% url 'users:register' %}">register</a> -
    <a href="{% url 'users:login' %}">log in</a>
{% endif %}
--snip--
```

Now users who are logged in see a personalized greeting and a logout link. Users not logged in see a registration page link and a login link. Try out the registration page by making several user accounts with different usernames.

In the next section, we'll restrict some of the pages so they're available only to registered users, and we'll make sure every topic belongs to a specific user.

### NOTE

*The registration system we've set up allows anyone to make any number of accounts for Learning Log. But some systems require users to confirm their identity by sending a confirmation email the user must reply to. By doing so, the system generates fewer spam accounts than the simple system we're using here. However, when you're learning to build apps, it's perfectly appropriate to practice with a simple user registration system like the one we're using.*

### TRY IT YOURSELF

**19-2. Blog Accounts:** Add a user authentication and registration system to the Blog project you started in Exercise 19-1 (page 438). Make sure logged-in users see their username somewhere on the screen and unregistered users see a link to the registration page.

## Allowing Users to Own Their Data

Users should be able to enter data exclusive to them, so we'll create a system to figure out which data belongs to which user, and then we'll restrict access to certain pages so users can work with only their own data.

In this section, we'll modify the Topic model so every topic belongs to a specific user. This will also take care of entries, because every entry belongs to a specific topic. We'll start by restricting access to certain pages.

## Restricting Access with @login\_required

Django makes it easy to restrict access to certain pages to logged-in users through the `@login_required` decorator. A *decorator* is a directive placed just before a function definition that Python applies to the function before it runs to alter how the function code behaves. Let's look at an example.

### Restricting Access to the Topics Page

Each topic will be owned by a user, so only registered users should be able to request the topics page. Add the following code to *learning\_logs/views.py*:

---

```
views.py  --snip--
from django.core.urlresolvers import reverse
from django.contrib.auth.decorators import login_required

from .models import Topic, Entry
--snip--

@login_required
def topics(request):
    """Show all topics."""
    --snip--
```

---

We first import the `login_required()` function. We apply `login_required()` as a decorator to the `topics()` view function by prepending `login_required` with the `@` symbol so Python knows to run the code in `login_required()` before the code in `topics()`.

The code in `login_required()` checks to see if a user is logged in, and Django will run the code in `topics()` only if they are. If the user is not logged in, they're redirected to the login page.

To make this redirect work, we need to modify *settings.py* so Django knows where to find the login page. Add the following at the very end of *settings.py*:

---

```
settings.py  """
Django settings for learning_log project
--snip--

# My settings
LOGIN_URL = '/users/login/'
```

---

Now when an unauthenticated user requests a page protected by the `@login_required` decorator, Django will send the user to the URL defined by `LOGIN_URL` in *settings.py*.

You can test this setting by logging out of any user accounts and going to the home page. Next, click the Topics link, which should redirect you to the login page. Then log in to any of your accounts, and from the home page click the Topics link again. You should be able to reach the topics page.

## Restricting Access Throughout Learning Log

Django makes it easy to restrict access to pages, but you have to decide which pages to protect. It's better to think about which pages need to be unrestricted first and then restrict all the other pages in the project. You can easily correct overrestricting access, and it's less dangerous than leaving sensitive pages unrestricted.

In Learning Log, we'll keep the home page, the registration page, and logout unrestricted. We'll restrict access to every other page.

Here's *learning\_logs/views.py* with `@login_required` decorators applied to every view except `index()`:

---

```
views.py  --snip--
@login_required
def topics(request):
    --snip--

@login_required
def topic(request, topic_id):
    --snip--

@login_required
def new_topic(request):
    --snip--

@login_required
def new_entry(request, topic_id):
    --snip--

@login_required
def edit_entry(request, entry_id):
    --snip--
```

---

Try accessing each of these pages while logged out: you'll be redirected back to the login page. You'll also be unable to click links to pages such as `new_topic`. But if you enter the URL `http://localhost:8000/new_topic/`, you'll be redirected to the login page. You should restrict access to any URL that's publicly accessible and relates to private user data.

## Connecting Data to Certain Users

Now we need to connect the data submitted to the user who submitted it. We need to connect only the data highest in the hierarchy to a user, and the lower-level data will follow. For example, in Learning Log, topics are the highest level of data in the app, and all entries are connected to a topic. As long as each topic belongs to a specific user, we'll be able to trace the ownership of each entry in the database.

We'll modify the `Topic` model by adding a foreign key relationship to a user. We'll then have to migrate the database. Finally, we'll have to modify some of the views so they only show the data associated with the currently logged-in user.

## Modifying the Topic Model

The modification to *models.py* is just two lines:

```
models.py
from django.db import models
from django.contrib.auth.models import User

class Topic(models.Model):
    """A topic the user is learning about"""
    text = models.CharField(max_length=200)
    date_added = models.DateTimeField(auto_now_add=True)
    owner = models.ForeignKey(User)

    def __str__(self):
        """Return a string representation of the model."""
        return self.text

class Entry(models.Model):
    --snip--
```

We first import the *User* model from *django.contrib.auth*. We then add an *owner* field to *Topic*, which establishes a foreign key relationship to the *User* model.

## Identifying Existing Users

When we migrate the database, Django will modify the database so it can store a connection between each topic and a user. To make the migration, Django needs to know which user to associate with each existing topic. The simplest approach is to give all existing topics to one user—for example, the superuser. First, we need to know the ID of that user.

Let's look at the IDs of all users created so far. Start a Django shell session and issue the following commands:

```
(venv)learning_log$ python manage.py shell
❶ >>> from django.contrib.auth.models import User
❷ >>> User.objects.all()
[<User: ll_admin>, <User: eric>, <User: willie>]
❸ >>> for user in User.objects.all():
...     print(user.username, user.id)
...
ll_admin 1
eric 2
willie 3
>>>
```

At ❶ we import the *User* model into the shell session. We then look at all the users that have been created so far ❷. The output shows three users: *ll\_admin*, *eric*, and *willie*.

At ❸ we loop through the list of users and print each user's username and ID. When Django asks which user to associate the existing topics with, we'll use one of these ID values.



## Migrating the Database

Now that we know the IDs, we can migrate the database.

---

```
❶ (venv)learning_log$ python manage.py makemigrations learning_logs
❷ You are trying to add a non-nullable field 'owner' to topic without a default;
  we can't do that (the database needs something to populate existing rows).
❸ Please select a fix:
    1) Provide a one-off default now (will be set on all existing rows)
    2) Quit, and let me add a default in models.py
❹ Select an option: 1
❺ Please enter the default value now, as valid Python
  The datetime and django.utils.timezone modules are available, so you can do
  e.g. timezone.now()
❻ >>> 1
  Migrations for 'learning_logs':
    0003_topic_owner.py:
      - Add field owner to topic
```

---

We start by issuing the `makemigrations` command ❶. In the output at ❷, Django indicates that we're trying to add a required (non-nullable) field to an existing model (`topic`) with no default value specified. Django gives us two options at ❸: we can provide a default right now, or we can quit and add a default value in `models.py`. At ❹ we've chosen the first option. Django then asks us to enter the default value ❺.

To associate all existing topics with the original admin user, `ll_admin`, I entered the user ID of 1 at ❻. You can use the ID of any user you've created; it doesn't have to be a superuser. Django then migrates the database using this value and generates the migration file `0003_topic_owner.py`, which adds the field `owner` to the `Topic` model.

Now we can carry out the migration. Enter the following in an active virtual environment:

---

```
(venv)learning_log$ python manage.py migrate
Operations to perform:
  Synchronize unmigrated apps: messages, staticfiles
  Apply all migrations: learning_logs, contenttypes, sessions, admin, auth
--snip--
Running migrations:
  Rendering model states... DONE
❶ Applying learning_logs.0003_topic_owner... OK
(venv)learning_log$
```

---

Django applies the new migration, and the result is OK ❶.

We can verify that the migration worked as expected in the shell session, like this:

---

```
❶ >>> from learning_logs.models import Topic
❷ >>> for topic in Topic.objects.all():
...     print(topic, topic.owner)
```

```
...
Chess ll_admin
Rock Climbing ll_admin
>>>
```

---

We import `Topic` from `learning_logs.models` ❶ and then loop through all existing topics, printing each topic and the user it belongs to ❷. You can see that each topic now belongs to the user `ll_admin`.

**NOTE**

*You can simply reset the database instead of migrating, but that will lose all existing data. It's good practice to learn how to migrate a database while maintaining the integrity of users' data. If you do want to start with a fresh database, issue the command `python manage.py flush` to rebuild the database structure. You'll have to create a new superuser, and all of your data will be gone.*

## Restricting Topics Access to Appropriate Users

Currently, if you're logged in, you'll be able to see all the topics, no matter which user you're logged in as. We'll change that by showing users only the topics that belong to them.

Make the following change to the `topics()` function in `views.py`:

`views.py`

```
--snip--
@login_required
def topics(request):
    """Show all topics."""
    topics = Topic.objects.filter(owner=request.user).order_by('date_added')
    context = {'topics': topics}
    return render(request, 'learning_logs/topics.html', context)
--snip--
```

---

When a user is logged in, the `request` object has a `request.user` attribute set that stores information about the user. The code fragment `Topic.objects.filter(owner=request.user)` tells Django to retrieve only the `Topic` objects from the database whose `owner` attribute matches the current user. Because we're not changing how the topics are displayed, we don't need to change the template for the topics page at all.

To see if this works, log in as the user you connected all existing topics to, and go to the topics page. You should see all the topics. Now log out, and log back in as a different user. The topics page should list no topics.

## Protecting a User's Topics

We haven't actually restricted access to the topic pages yet, so any registered user could try a bunch of URLs, like `http://localhost:8000/topics/1/`, and retrieve topic pages that happen to match.

Try it yourself. While logged in as the user that owns all topics, copy the URL or note the ID in the URL of a topic, and then log out and log back in

as a different user. Enter the URL of that topic. You should be able to read the entries, even though you're logged in as a different user.

We'll fix this now by performing a check before retrieving the requested entries in the `topic()` view function:

---

```
views.py  from django.shortcuts import render
          ❶ from django.http import HttpResponseRedirect, Http404
          from django.core.urlresolvers import reverse
          --snip--

          @login_required
          def topic(request, topic_id):
              """Show a single topic and all its entries."""
              topic = Topic.objects.get(id=topic_id)
              # Make sure the topic belongs to the current user.
          ❷  if topic.owner != request.user:
              raise Http404

              entries = topic.entry_set.order_by('-date_added')
              context = {'topic': topic, 'entries': entries}
              return render(request, 'learning_logs/topic.html', context)
          --snip--
```

---

A 404 response is a standard error response that's returned when a requested resource doesn't exist on a server. Here we import the `Http404` exception ❶, which we'll raise if the user requests a topic they shouldn't see. After receiving a topic request, we make sure the topic's user matches the currently logged-in user before rendering the page. If the current user doesn't own the requested topic, we raise the `Http404` exception ❷, and Django returns a 404 error page.

Now if you try to view another user's topic entries, you'll see a *Page Not Found* message from Django. In Chapter 20, we'll configure the project so users will see a proper error page.

### ***Protecting the edit\_entry Page***

The `edit_entry` pages have URLs in the form `http://localhost:8000/edit_entry/entry_id/`, where the `entry_id` is a number. Let's protect this page so no one can use the URL to gain access to someone else's entries:

---

```
views.py  --snip--
          @login_required
          def edit_entry(request, entry_id):
              """Edit an existing entry."""
              entry = Entry.objects.get(id=entry_id)
              topic = entry.topic
              if topic.owner != request.user:
                  raise Http404
```

```
if request.method != 'POST':
    # Initial request; pre-fill form with the current entry.
    --snip--
```

---

We retrieve the entry and the topic associated with this entry. We then check if the owner of the topic matches the currently logged-in user; if they don't match, we raise an `Http404` exception.

### ***Associating New Topics with the Current User***

Currently, our page for adding new topics is broken, because it doesn't associate new topics with any particular user. If you try adding a new topic, you'll see the error message `IntegrityError` along with `learning_logs_topic.user_id` may not be `NULL`. Django's saying you can't create a new topic without specifying a value for the topic's owner field.

There's a straightforward fix for this problem, because we have access to the current user through the request object. Add the following code, which associates the new topic with the current user:

---

```
views.py    --snip--
@login_required
def new_topic(request):
    """Add a new topic."""
    if request.method != 'POST':
        # No data submitted; create a blank form.
        form = TopicForm()
    else:
        # POST data submitted; process data.
        form = TopicForm(request.POST)
        if form.is_valid():
            ❶ new_topic = form.save(commit=False)
            ❷ new_topic.owner = request.user
            ❸ new_topic.save()
            return HttpResponseRedirect(reverse('learning_logs:topics'))

    context = {'form': form}
    return render(request, 'learning_logs/new_topic.html', context)
--snip--
```

---

When we first call `form.save()`, we pass the `commit=False` argument because we need to modify the new topic before saving it to the database ❶. We then set the new topic's `owner` attribute to the current user ❷. Finally, we call `save()` on the topic instance just defined ❸. Now the topic has all the required data and will save successfully.

You should be able to add as many new topics as you want for as many different users as you want. Each user will have access only to their own data, whether they're viewing data, entering new data, or modifying old data.

### TRY IT YOURSELF

**19-3. Refactoring:** There are two places in *views.py* where we make sure the user associated with a topic matches the currently logged-in user. Put the code for this check in a function called `check_topic_owner()`, and call this function where appropriate.

**19-4. Protecting new\_entry:** A user can add a new entry to another user's learning log by entering a URL with the ID of a topic belonging to another user. Prevent this attack by checking that the current user owns the entry's topic before saving the new entry.

**19-5. Protected Blog:** In your Blog project, make sure each blog post is connected to a particular user. Make sure all posts are publicly accessible but only registered users can add posts and edit existing posts. In the view that allows users to edit their posts, make sure the user is editing their own post before processing the form.

## Summary

In this chapter you learned to use forms to allow users to add new topics and entries, and edit existing entries. You then learned how to implement user accounts. You allowed existing users to log in and out, and you learned how to use Django's default `UserCreationForm` to let people create new accounts.

After building a simple user authentication and registration system, you restricted access to logged-in users for certain pages using the `@login_required` decorator. You then attributed data to specific users through a foreign key relationship. You also learned to migrate the database when the migration requires you to specify some default data.

Finally, you learned how to make sure a user can see only data that belongs to them by modifying the view functions. You retrieved appropriate data using the `filter()` method, and you learned to compare the owner of the requested data to the currently logged-in user.

It may not always be immediately obvious what data you should make available and what data you should protect, but this skill will come with practice. The decisions we've made in this chapter to secure our users' data illustrate why working with others is a good idea when building a project: having someone else look over your project makes it more likely that you'll spot vulnerable areas.

We now have a fully functioning project running on our local machine. In the final chapter we'll style Learning Log to make it visually appealing, and we'll deploy the project to a server so anyone with Internet access can register and make an account.

# 20

## STYLING AND DEPLOYING AN APP



Learning Log is fully functional now, but it has no styling and runs only on your local machine. In this chapter we'll style the project in a simple but professional manner and then deploy it to a live server so anyone in the world can make an account.

For the styling we'll use the Bootstrap library, a collection of tools for styling web applications so they look professional on all modern devices, from a large flat-screen monitor to a smartphone. To do this, we'll use the `django-bootstrap3` app, which will also give you practice using apps made by other Django developers.

We'll deploy Learning Log using Heroku, a site that lets you push your project to one of its servers, making it available to anyone with an Internet connection. We'll also start using a version control system called Git to track changes to the project.

When you're finished with Learning Log, you'll be able to develop simple web applications, make them look good, and deploy them to a live server. You'll also be able to use more advanced learning resources as you develop your skills.

## Styling Learning Log

We've purposely ignored styling until now to focus on Learning Log's functionality first. This is a good way to approach development, because an app is useful only if it works. Of course, once it's working, appearance is critical so people will want to use it.

In this section I'll introduce the `django-bootstrap3` app and show you how to integrate it into a project to make it ready for live deployment.

### *The django-bootstrap3 App*

We'll use `django-bootstrap3` to integrate Bootstrap into our project. This app downloads the required Bootstrap files, places them in an appropriate location in your project, and makes the styling directives available in your project's templates.

To install `django-bootstrap3`, issue the following command in an active virtual environment:

---

```
(ll_env)learning_log$ pip install django-bootstrap3
--snip--
Successfully installed django-bootstrap3
```

---

Next, we need to add the following code to include `django-bootstrap3` in `INSTALLED_APPS` in `settings.py`:

```
settings.py  --snip--
INSTALLED_APPS = (
    --snip--
    'django.contrib.staticfiles',

    # Third party apps
    'bootstrap3',

    # My apps
    'learning_logs',
    'users',
)
--snip--
```

---

Start a new section called *Third party apps* for apps created by other developers and add `'bootstrap3'` to this section. Most apps need to be included in `INSTALLED_APPS`, but to be sure, read the setup instructions for the particular app you're using.

We need `django-bootstrap3` to include jQuery, a JavaScript library that enables some of the interactive elements that the Bootstrap template provides. Add this code to the end of `settings.py`:

```
settings.py  --snip--
# My settings
LOGIN_URL = '/users/login/'

# Settings for django-bootstrap3
BOOTSTRAP3 = {
    'include_jquery': True,
}
```

This code spares us from having to download jQuery and place it in the correct location manually.

## Using Bootstrap to Style Learning Log

Bootstrap is basically a large collection of styling tools. It also has a number of templates you can apply to your project to create a particular overall style. If you're just starting out, it's much easier to use these templates than it is to use individual styling tools. To see the templates Bootstrap offers, go to the *Getting Started* section at <http://getbootstrap.com/>; then scroll down to the *Examples* heading, and look for the *Navbars in action* section. We'll use the *Static top navbar* template, which provides a simple top navigation bar, a page header, and a container for the content of the page.

Figure 20-1 shows what the home page will look like after we apply Bootstrap's template to `base.html` and modify `index.html` slightly.




Figure 20-1: The Learning Log home page using Bootstrap



Now that you know the result we're after, the sections that follow will be easier to understand.

## Modifying *base.html*

We need to modify the *base.html* template to accommodate the Bootstrap template. I'll introduce the new *base.html* in parts.

### Defining the HTML Headers

The first change to *base.html* defines the HTML headers in the file so whenever a Learning Log page is open, the browser title bar displays the site name. We'll also add some requirements for using Bootstrap in our templates. Delete everything in *base.html* and replace it with the following code:

---

```
base.html ❶ {% load bootstrap3 %}

❷ <!DOCTYPE html>
❸ <html lang="en">
❹ <head>
    <meta charset="utf-8">
    <meta http-equiv="X-UA-Compatible" content="IE=edge">
    <meta name="viewport" content="width=device-width, initial-scale=1">

❺ <title>Learning Log</title>

❻ {% bootstrap_css %}
    {% bootstrap_javascript %}

❼ </head>
```

---

At ❶ we load the collection of template tags available in `django-bootstrap3`. Next, we declare this file as an HTML document ❷ written in English ❸. An HTML file is divided into two main parts, the *head* and the *body*—the head of the file begins at ❹. The head of an HTML file doesn't contain any content: it just tells the browser what it needs to know to display the page correctly. At ❺ we include a title element for the page, which will be displayed in the title bar of the browser whenever Learning Log is open.

At ❻ we use one of `django-bootstrap3`'s custom template tags, which tells Django to include all the Bootstrap style files. The tag that follows enables all the interactive behavior you might use on a page, such as collapsible navigation bars. At ❼ is the closing `</head>` tag.

### Defining the Navigation Bar

Now we'll define the navigation bar at the top of the page:

---

```
--snip--
</head>

<body>

    <!-- Static navbar -->
```

```

❶ <nav class="navbar navbar-default navbar-static-top">
    <div class="container">

        <div class="navbar-header">
❷     <button type="button" class="navbar-toggle collapsed"
        data-toggle="collapse" data-target="#navbar"
        aria-expanded="false" aria-controls="navbar">
        </button>
❸     <a class="navbar-brand" href="{% url 'learning_logs:index' %}">
        Learning Log</a>
    </div>

❹    <div id="navbar" class="navbar-collapse collapse">
❺    <ul class="nav navbar-nav">
❻    <li><a href="{% url 'learning_logs:topics' %}">Topics</a></li>
    </ul>

❽    <ul class="nav navbar-nav navbar-right">
        {% if user.is_authenticated %}
        <li><a>Hello, {{ user.username }}.</a></li>
        <li><a href="{% url 'users:logout' %}">log out</a></li>
        {% else %}
        <li><a href="{% url 'users:register' %}">register</a></li>
        <li><a href="{% url 'users:login' %}">log in</a></li>
        {% endif %}
❾    </ul>
    </div><!--/.nav-collapse -->

    </div>
</nav>

```

The first element is the opening `<body>` tag. The body of an HTML file contains the content users will see on a page. At ❶ is a `<nav>` element that indicates the navigation links section of the page. Everything contained in this element is styled according to the Bootstrap style rules defined by the selectors `navbar`, `navbar-default`, and `navbar-static-top`. A *selector* determines which elements on a page a certain style rule applies to.

At ❷ the template defines a button that will appear if the browser window is too narrow to display the whole navigation bar horizontally. When the user clicks the button, the navigation elements will appear in a drop-down list. The `collapse` reference causes the navigation bar to collapse when the user shrinks the browser window or when the site is displayed on mobile devices with small screens.

At ❸ we set the project's name to appear at the far left of the navigation bar and make it a link to the home page, because it will appear on every page in the project.

At ❹ we define a set of links that lets users navigate the site. A navigation bar is basically a list that starts with `<ul>` ❺, and each link is an item in this list (`<li>`) ❻. To add more links, insert more lines using the following structure:

---

```
<li><a href="{% url 'learning_logs:title' %}">Title</a></li>
```

---

This line represents a single link in the navigation bar. The link is taken directly from the previous version of *base.html*.

At ⑦ we place a second list of navigation links, this time using the selector `navbar-right`. The `navbar-right` selector styles the set of links so it appears at the right edge of the navigation bar where you typically see login and registration links. Here we'll display the user greeting and logout link or links to register or log in. The rest of the code in this section closes out the elements that contain the navigation bar ⑧.

## Defining the Main Part of the Page

The rest of *base.html* contains the main part of the page:

---

```
--snip--
    </nav>

❶    <div class="container">

        <div class="page-header">
❷        {% block header %}{% endblock header %}
        </div>
        <div>
❸        {% block content %}{% endblock content %}
        </div>

    </div> <!-- /container -->

</body>
</html>
```

---

At ❶ is an opening `div` with the class `container`. A *div* is a section of a web page that can be used for any purpose and can be styled with a border, space around the element (margins), space between the contents and the border (padding), background colors, and other style rules. This particular `div` acts as a container into which we place two elements: a new block called `header` ❷ and the content block we used in Chapter 18 ❸. The header block contains information telling the user what kind of information the page holds and what they can do on a page. It has the class `page-header`, which applies a set of style rules to the block. The content block is in a separate `div` with no specific style classes.

When you load the home page of Learning Log in a browser, you should see a professional-looking navigation bar that matches the one shown in Figure 20-1. Try resizing the window so it's really narrow; the navigation bar should be replaced by a button. Click the button, and all the links should appear in a drop-down list.

### NOTE

*This simplified version of the Bootstrap template should work on most recent browsers. Earlier browsers may not render some styles correctly. The full template, available at <http://getbootstrap.com/getting-started/#examples/>, will work on almost all available browsers.*

## Styling the Home Page Using a Jumbotron

Let's update the home page using the newly defined header block and another Bootstrap element called a *jumbotron*—a large box that will stand out from the rest of the page and can contain anything you want. It's typically used on home pages to hold a brief description of the overall project. While we're at it, we'll update the message on the home page as well. Here's *index.html*:

---

```
index.html    {% extends "learning_logs/base.html" %}

❶ {% block header %}
❷   <div class='jumbotron'>
      <h1>Track your learning.</h1>
    </div>
  {% endblock header %}

  {% block content %}
❸   <h2>
      <a href="{% url 'users:register' %}">Register an account</a> to make
      your own Learning Log, and list the topics you're learning about.
    </h2>
    <h2>
      Whenever you learn something new about a topic, make an entry
      summarizing what you've learned.
    </h2>
  {% endblock content %}
```

---

At ❶ we tell Django that we're about to define what goes in the header block. Inside a *jumbotron* element ❷ we place a short tagline, *Track your learning*, to give first-time visitors a sense of what Learning Log does.

At ❸ we add text to provide a little more direction. We invite people to make an account, and we describe the two main actions—add new topics and make topic entries. The index page now looks like Figure 20-1 and is a significant improvement over our unstyled project.

## Styling the Login Page

We've refined the overall appearance of the login page but not the login form yet, so let's make the form look consistent with the rest of the page:

---

```
login.html    {% extends "learning_logs/base.html" %}

❶ {% load bootstrap3 %}

❷ {% block header %}
      <h2>Log in to your account.</h2>
  {% endblock header %}

  {% block content %}

❸   <form method="post" action="{% url 'users:login' %}" class="form">
      {% csrf_token %}
```

```

❹    {% bootstrap_form form %}

❺    {% buttons %}
      <button name="submit" class="btn btn-primary">log in</button>
    {% endbuttons %}

    <input type="hidden" name="next" value="{% url 'learning_logs:index' %}" />
  </form>

{% endblock content %}

```

---

At ❹ we load the bootstrap3 template tags into this template. At ❺ we define the header block, which describes what the page is for. Notice that we've removed the `{% if form.errors %}` block from the template; django-bootstrap3 manages form errors automatically.

At ❸ we add a `class="form"` attribute, and then we use the template tag `{% bootstrap_form %}` when we display the form ❹; this replaces the `{{ form.as_p }}` tag we were using in Chapter 19. The `{% bootstrap_form %}` template tag inserts Bootstrap style rules into the individual elements of the form as it's rendered. At ❺ we open a bootstrap3 template tag `{% buttons %}`, which adds Bootstrap styling to buttons.

Figure 20-2 shows the login form as it's rendered now. The page is much cleaner and has consistent styling and a clear purpose. Try logging in with an incorrect username or password; you'll see that even the error messages are styled consistently and integrate well with the overall site.




Figure 20-2: The login page styled with Bootstrap

## Styling the new\_topic Page

Let's make the rest of the pages look consistent as well. We'll update the new\_topic page next:

```
new_topic.html  {% extends "learning_logs/base.html" %}
                {% load bootstrap3 %}

                ❶ {% block header %}
                  <h2>Add a new topic:</h2>
                {% endblock header %}

                {% block content %}

                ❷ <form action="{% url 'learning_logs:new_topic' %}" method='post'
                  class="form">

                  {% csrf_token %}
                ❸ {% bootstrap_form form %}

                ❹ {% buttons %}
                  <button name="submit" class="btn btn-primary">add topic</button>
                {% endbuttons %}

                </form>

                {% endblock content %}
```

Most of the changes here are similar to those applied in *login.html*: we load bootstrap3 and add the header block with an appropriate message at ❶. Then we add the class="form" attribute to the <form> tag ❷, use the {% bootstrap\_form %} template tag instead of {{ form.as\_p }} ❸, and use the bootstrap3 structure for the submit button ❹. Log in and navigate to the new\_topic page; it should look similar to the login page now.

## Styling the Topics Page

Now let's make sure the pages for viewing information are styled appropriately as well, starting with the topics page:

```
topics.html    {% extends "learning_logs/base.html" %}

                ❶ {% block header %}
                  <h1>Topics</h1>
                {% endblock header %}

                {% block content %}

                <ul>
                  {% for topic in topics %}
                    <li>
```

```

❷      <h3>
        <a href="{% url 'learning_logs:topic' topic.id %}">{{ topic }}</a>
      </h3>
    </li>
    {% empty %}
    <li>No topics have been added yet.</li>
    {% endfor %}
  </ul>

❸  <h3><a href="{% url 'learning_logs:new_topic' %}">Add new topic</h3>

    {% endblock content %}

```

---

We don't need the `{% load bootstrap3 %}` tag, because we're not using any custom bootstrap3 template tags in this file. We add the heading *Topics* inside the header block ❶. We style each topic as an `<h3>` element to make them a little larger on the page ❷ and do the same for the link to add a new topic ❸.

## Styling the Entries on the Topic Page

The topic page has more content than most pages, so it needs a bit more work. We'll use Bootstrap's panels to make each entry stand out. A *panel* is a div with predefined styling and is perfect for displaying a topic's entries:

```

topic.html    {% extends 'learning_logs/base.html' %}

❶  {% block header %}
    <h2>{{ topic }}</h2>
  {% endblock header %}

  {% block content %}
    <p>
      <a href="{% url 'learning_logs:new_entry' topic.id %}">add new entry</a>
    </p>

    {% for entry in entries %}
❷    <div class="panel panel-default">
❸      <div class="panel-heading">
❹        <h3>
          {{ entry.date_added|date:'M d, Y H:i' }}
❺        <small>
          <a href="{% url 'learning_logs:edit_entry' entry.id %}">
            edit entry</a>
          </small>
        </h3>
      </div>
❻    <div class="panel-body">
      {{ entry.text|linebreaks }}
    </div>
  </div> <!-- panel -->
    {% endfor %}
  {% endblock content %}

```

```
{% empty %}
    There are no entries for this topic yet.
{% endfor %}
```

```
{% endblock content %}
```

We first place the topic in the header block ❶. We then delete the unordered list structure previously used in this template. Instead of making each entry a list item, we create a panel div element at ❷, which contains two more nested divs: a panel-heading div ❸ and a panel-body div ❹. The panel-heading div contains the date for the entry and the link to edit the entry. Both are styled as `<h3>` elements ❺, but we add `<small>` tags around the `edit_entry` link to make it a little smaller than the timestamp ❻.

At ❹ is the panel-body div, which contains the actual text of the entry. Notice that the Django code for including the information on the page hasn't changed at all; only the elements that affect the appearance of the page have changed.

Figure 20-3 shows the topic page with its new look. The functionality of Learning Log hasn't changed, but it looks more professional and inviting to users.

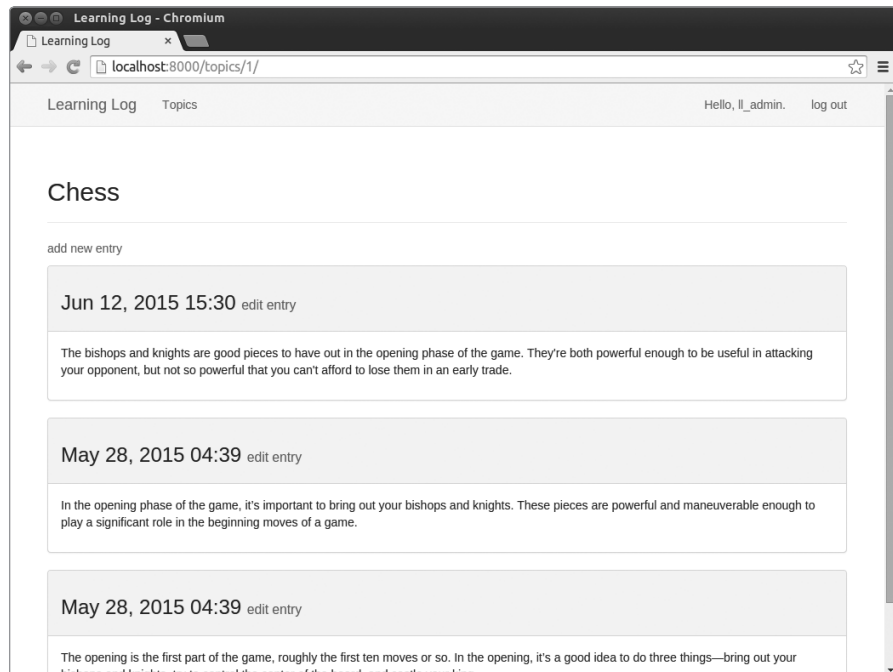


Figure 20-3: The topic page with Bootstrap styling

#### NOTE

*If you want to use a different Bootstrap template, follow a similar process to what we've done so far in this chapter. Copy the template into `base.html`, and modify the elements that contain actual content so the template displays your project's information. Then use Bootstrap's individual styling tools to style the content on each page.*



### TRY IT YOURSELF

**20-1. Other Forms:** We've applied Bootstrap's styles to the login and add\_topic pages. Make similar changes to the rest of the form-based pages: new\_entry and edit\_entry, and register.

**20-2. Stylish Blog:** Use Bootstrap to style the Blog project you created in Chapter 19.

## Deploying Learning Log

Now that we have a professional-looking project, let's deploy it to a live server so anyone with an internet connection can use it. We'll use Heroku, a web-based platform that allows you to manage the deployment of web applications. We'll get Learning Log up and running on Heroku.

The process is slightly different on Windows than it is on Linux and OS X. If you're using Windows, check for notes in each section that specify what you'll need to do differently on your system.

### ***Making a Heroku Account***

To make an account, go to <https://heroku.com/> and click one of the signup links. It's free to make an account, and Heroku has a free tier that allows you to test your projects in live deployment.

#### **NOTE**

*Heroku's free tier has limits, such as the number of apps you can deploy and how often people can visit your app. But these limits are generous enough to let you practice deploying apps without any cost.*

### ***Installing the Heroku Toolbelt***

To deploy and manage a project on Heroku's servers, you'll need the tools available in the Heroku Toolbelt. To install the latest version of the Heroku Toolbelt, visit <https://toolbelt.heroku.com/> and follow the directions for your operating system, which will include either a one-line terminal command or an installer you can download and run.

### ***Installing Required Packages***

You'll also need to install a number of packages that help serve Django projects on a live server. In an active virtual environment, issue the following commands:

---

```
(ll_env)learning_log$ pip install dj-database-url
(ll_env)learning_log$ pip install dj-static
(ll_env)learning_log$ pip install static3
(ll_env)learning_log$ pip install gunicorn
```

---

Make sure you issue the commands one at a time so you know if any package fails to install correctly. The package `dj-database-url` helps Django communicate with the database Heroku uses, `dj-static` and `static3` help Django manage static files correctly, and `gunicorn` is a server capable of serving apps in a live environment. (*Static files* contain style rules and JavaScript files.)

**NOTE**

*Some of the required packages may not install on Windows, so don't be concerned if you get an error message when you try to install some of them. What matters is getting Heroku to install the packages on the live deployment, and we'll do that in the next section.*

### Creating a Packages List with a `requirements.txt` File

Heroku needs to know which packages our project depends on, so we'll use `pip` to generate a file listing them. Again, from an active virtual environment, issue the following command:

---

```
(ll_env)learning_log$ pip freeze > requirements.txt
```

---

The `freeze` command tells `pip` to write the names of all the packages currently installed in the project into the file `requirements.txt`. Open `requirements.txt` to see the packages and version numbers installed in your project (Windows users might not see all of these lines):

---

```
requirements.txt Django==1.8.4
dj-database-url==0.3.0
dj-static==0.0.6
django-bootstrap3==6.2.2
gunicorn==19.3.0
static3==0.6.1
```

---

Learning Log already depends on six different packages with specific version numbers, so it requires a specific environment to run properly. When we deploy Learning Log, Heroku will install all the packages listed in `requirements.txt`, creating an environment with the same packages we're using locally. For this reason, we can be confident the deployed project will behave the same as it does on our local system. This is a huge advantage as you start to build and maintain various projects on your system.

Next, we need to add `psycopg2`, which helps Heroku manage the live database, to the list of packages. Open `requirements.txt` and add the line `psycopg2>=2.6.1`. This will install version 2.6.1 of `psycopg2`, or a newer version if it's available:

---

```
requirements.txt Django==1.8.4
dj-database-url==0.3.0
dj-static==0.0.6
django-bootstrap3==6.2.2
gunicorn==19.3.0
psycopg2>=2.6.1
```

---

```
static3==0.6.1
psycopg2>=2.6.1
```

---

If any of the packages didn't install on your system, add those as well. When you're finished, your *requirements.txt* file should include each of the packages shown above. If a package is listed on your system but the version number differs from what's shown here, keep the version you have on your system.

**NOTE**

*If you're using Windows, make sure your version of requirements.txt matches the list shown here regardless of which packages you were able to install on your system.*

## Specifying the Python Runtime

Unless you specify a Python version, Heroku will use its own current default version of Python. Let's make sure Heroku uses the same version of Python we're using. In an active virtual environment, issue the command `python --version`:

---

```
(ll_env)learning_log$ python --version
Python 3.5.0
```

---

In this example I'm running Python 3.5.0. Make a new file called *runtime.txt* in the same directory as *manage.py*, and enter the following:

```
runtime.txt  python-3.5.0
```

---

This file should contain one line with your Python version specified in the format shown; make sure you enter `python` in lowercase, followed by a hyphen, followed by the three-part version number.

**NOTE**

*If you get an error reporting that the Python runtime you requested is not available, go to <https://devcenter.heroku.com/> and click **Python**; then look for a link to Specifying a Python Runtime. Scan through the article to find the available runtimes, and use the one that most closely matches your Python version.*

## Modifying settings.py for Heroku

Now we need to add a section at the end of *settings.py* to define some settings specifically for the Heroku environment:

```
settings.py  --snip--
# Settings for django-bootstrap3
BOOTSTRAP3 = {
    'include_jquery': True,
}
```

```

# Heroku settings
❶ if os.getcwd() == '/app':
❷     import dj_database_url
        DATABASES = {
            'default': dj_database_url.config(default='postgres://localhost')
        }

        # Honor the 'X-Forwarded-Proto' header for request.is_secure().
❸     SECURE_PROXY_SSL_HEADER = ('HTTP_X_FORWARDED_PROTO', 'https')

        # Allow all host headers.
❹     ALLOWED_HOSTS = ['*']

        # Static asset configuration
❺     BASE_DIR = os.path.dirname(os.path.abspath(__file__))
        STATIC_ROOT = 'staticfiles'
        STATICFILES_DIRS = (
            os.path.join(BASE_DIR, 'static'),
        )

```

---

At ❶ we use the function `getcwd()`, which gets the *current working directory* the file is running from. In a Heroku deployment, the directory is always `/app`. In a local deployment, the directory is usually the name of the project folder (*learning\_log* in our case). The `if` test ensures that the settings in this block apply only when the project is deployed on Heroku. This structure allows us to have one settings file that works for our local development environment as well as the live server.

At ❷ we import `dj_database_url` to help configure the database on Heroku. Heroku uses PostgreSQL (also called Postgres), a more advanced database than SQLite, and these settings configure the project to use Postgres on Heroku. The rest of the settings support HTTPS requests ❸, ensure that Django will serve the project from Heroku's URL ❹, and set up the project to serve static files correctly on Heroku ❺.

## Making a Procfile to Start Processes

A *Procfile* tells Heroku which processes to start in order to serve the project properly. This is a one-line file that you should save as *Procfile*, with an uppercase *P* and no file extension, in the same directory as *manage.py*.

Here's what goes in *Procfile*:

---

```
Procfile  web: gunicorn learning_log.wsgi --log-file -
```

---

This line tells Heroku to use gunicorn as a server and to use the settings in *learning\_log/wsgi.py* to launch the app. The `log-file` flag tells Heroku the kinds of events to log.

## Modifying wsgi.py for Heroku

We also need to modify *wsgi.py* for Heroku, because Heroku needs a slightly different setup than what we've been using:

---

```
wsgi.py  --snip--
import os

from django.core.wsgi import get_wsgi_application
from dj_static import Cling

os.environ.setdefault("DJANGO_SETTINGS_MODULE", "learning_log.settings")
application = Cling(get_wsgi_application())
```

---

We import *Cling*, which helps serve static files correctly, and use it to launch the application. This code will work locally as well, so we don't need to put it in an *if* block.

## Making a Directory for Static Files

On Heroku, Django collects all the static files and places them in one place so it can manage them efficiently. We'll create a directory for these static files. Inside the *learning\_log* folder we've been working from is another folder called *learning\_log*. In this nested folder, make a new folder called *static* with the path *learning\_log/learning\_log/static/*. We also need to make a placeholder file to store in this directory for now, because empty directories won't be included in the project when it's pushed to Heroku. In the *static/* directory, make a file called *placeholder.txt*:

---

```
placeholder.txt  This file ensures that learning_log/static/ will be added to the project.
                  Django will collect static files and place them in learning_log/static/.
```

---

There's nothing special about this text; it just reminds us why we included this file in the project.

## Using the gunicorn Server Locally

If you're using Linux or OS X, you can try using the gunicorn server locally before deploying to Heroku. From an active virtual environment, run the command *heroku local* to start the processes defined in *Procfile*:

---

```
(ll_env)learning_log$ heroku local
Installing Heroku Toolbelt v4... done
--snip--
forego | starting web.1 on port 5000
❶ web.1 | [2015-08-13 22:00:45 -0800] [12875] [INFO] Starting gunicorn 19.3.0
❷ web.1 | [2015-08-13 22:00:45 -0800] [12875] [INFO] Listening at:
        http://0.0.0.0:5000 (12875)
❸ web.1 | [2015-08-13 22:00:45 -0800] [12878] [INFO] Booting worker with pid: 12878
```

---

The first time you run `heroku local`, a number of packages from the Heroku Toolbelt will be installed. The output shows that gunicorn has been started with a process id of 12875 in this example ❶. At ❷ gunicorn is listening for requests on port 5000. In addition, gunicorn has started a *worker* process (12878) to help it serve requests ❸.

Visit <http://localhost:5000/> to make sure everything is working; you should see the Learning Log home page, just as it appears when you use the Django server (`runserver`). Press CTRL-C to stop the processes started by `heroku local`. You should continue to use `runserver` for local development.

**NOTE**

*gunicorn won't run on Windows, so skip this step if you're using Windows. This won't affect your ability to deploy the project to Heroku.*

## Using Git to Track the Project's Files

If you completed Chapter 17, you'll know that Git is a version control program that allows you to take a snapshot of the code in your project each time you implement a new feature successfully. This allows you to easily return to the last working snapshot of your project if anything goes wrong; for example, if you accidentally introduce a bug while working on a new feature. Each of these snapshots is called a *commit*.

Using Git means you can try implementing new features without worrying about breaking your project. When you're deploying to a live server, you need to make sure you're deploying a working version of your project. If you want to read more about Git and version control, see Appendix D.

## Installing Git

The Heroku Toolbelt includes Git, so it should already be installed on your system. But terminal windows that were open before you installed the Heroku Toolbelt won't have access to Git, so open a new terminal window and issue the command `git --version`:

---

```
(ll_env)learning_log$ git --version
git version 2.5.0
```

---

If you get an error message for some reason, see the instructions in Appendix D for installing Git.

## Configuring Git

Git keeps track of who makes changes to a project, even in cases like this when there's only one person working on the project. To do this, Git needs to know your username and email. You have to provide a username, but feel free to make up an email for your practice projects:

---

```
(ll_env)learning_log$ git config --global user.name "ehmatthes"
(ll_env)learning_log$ git config --global user.email "eric@example.com"
```

---

If you forget this step, Git will prompt you for this information when you make your first commit.

## Ignoring Files

We don't need Git to track every file in the project, so we'll tell Git to ignore some files. Make a file called `.gitignore` in the folder that contains `manage.py`. Notice that this filename begins with a dot and has no file extension. Here's what goes in `.gitignore`:

---

```
.gitignore ll_env/  
__pycache__/  
*.sqlite3
```

---

We tell Git to ignore the entire directory `ll_env`, because we can re-create it automatically at any time. We also don't track the `__pycache__` directory, which contains the `.pyc` files that are created automatically when Django runs the `.py` files. We don't track changes to the local database, because it's a bad habit: if you're ever using SQLite on a server, you might accidentally overwrite the live database with your local test database when you push the project to the server.

### NOTE

*If you're using Python 2.7, replace `__pycache__` with `*.pyc` because Python 2.7 doesn't create a `__pycache__` directory.*

## Committing the Project

We need to initialize a Git repository for Learning Log, add all the necessary files to the repository, and commit the initial state of the project. Here's how we do that:

---

```
❶ (ll_env)learning_log$ git init  
Initialized empty Git repository in /home/ehmatthes/pcc/learning_log/.git/  
❷ (ll_env)learning_log$ git add .  
❸ (ll_env)learning_log$ git commit -am "Ready for deployment to heroku."  
[master (root-commit) dbc1d99] Ready for deployment to heroku.  
 43 files changed, 746 insertions(+)  
 create mode 100644 .gitignore  
 create mode 100644 Procfile  
 --snip--  
 create mode 100644 users/views.py  
❹ (ll_env)learning_log$ git status  
# On branch master  
nothing to commit, working directory clean  
(ll_env)learning_log$
```

---

At ❶ we issue the `git init` command to initialize an empty repository in the directory containing Learning Log. At ❷ we use the `git add .` command, which adds all the files that aren't being ignored to the repository. (Don't forget the dot.) At ❸ we issue the command `git commit -am commit message`: the `-a` flag tells Git to include all changed files in this commit, and the `-m` flag tells Git to record a log message.

Issuing the `git status` command ❹ indicates that we're on the *master* branch and that our working directory is *clean*. This is the status you'll want to see any time you push your project to Heroku.

## Pushing to Heroku

We're finally ready to push the project to Heroku. In an active terminal session, issue the following commands:

---

```
❶ (ll_env)learning_log$ heroku login
Enter your Heroku credentials.
Email: eric@example.com
Password (typing will be hidden):
Logged in as eric@example.com
❷ (ll_env)learning_log$ heroku create
Creating afternoon-meadow-2775... done, stack is cedar-14
https://afternoon-meadow-2775.herokuapp.com/ |
https://git.heroku.com/afternoon-meadow-2775.git
Git remote heroku added
❸ (ll_env)learning_log$ git push heroku master
--snip--
remote: -----> Launching... done, v6
❹ remote: https://afternoon-meadow-2775.herokuapp.com/ deployed to Heroku
remote: Verifying deploy.... done.
To https://git.heroku.com/afternoon-meadow-2775.git
bdb2a35..62d711d master -> master
(ll_env)learning_log$
```

---

First, log in to Heroku in the terminal session with the username and password you used to create an account at <https://heroku.com/> ❶. Then tell Heroku to build an empty project ❷. Heroku generates a name made up of two words and a number; you can change this later on. We then issue the command `git push heroku master` ❸, which tells Git to push the master branch of the project to the repository Heroku just created. Heroku then builds the project on its servers using these files. At ❹ is the URL we'll use to access the live project.

When you've issued these commands, the project is deployed but not fully configured. To check that the server process started correctly, use the `heroku ps` command:

---

```
(ll_env)learning_log$ heroku ps
❶ Free quota left: 17h 40m
❷ === web (Free): `unicorn learning_log.wsgi __log-file -`
web.1: up 2015/08/14 07:08:51 (~ 10m ago)
(ll_env)learning_log$
```

---

The output shows how much more time the project can be active in the next 24 hours ❶. At the time of this writing, Heroku allows free deployments to be active for up to 18 hours in any 24-hour period. If a project



exceeds these limits, a standard server error page will be displayed; we'll customize this error page shortly. At ❷ we see that the process defined in *Procfile* has been started.

Now we can open the app in a browser using the command `heroku open`:

---

```
(ll_env)learning_log$ heroku open
Opening afternoon-meadow-2775... done
```

---

This command spares you from opening a browser and entering the URL Heroku showed you, but that's another way to open the site. You should see the home page for Learning Log, styled correctly. However, you can't use the app yet because we haven't set up the database.

**NOTE**

*Heroku's deployment process changes from time to time. If you have any issues you can't resolve, look at Heroku's documentation for help. Go to <https://devcenter.heroku.com/>, click **Python**, and look for a link to Getting Started with Django. If you can't understand what you see there, check out the suggestions in Appendix C.*

## Setting Up the Database on Heroku

We need to run `migrate` once to set up the live database and apply all the migrations we generated during development. You can run Django and Python commands on a Heroku project using the command `heroku run`. Here's how to run `migrate` on the Heroku deployment:

---

```
❶ (ll_env)learning_log$ heroku run python manage.py migrate
❷ Running `python manage.py migrate` on afternoon-meadow-2775... up, run.2435
--snip--
❸ Running migrations:
--snip--
Applying learning_logs.0001_initial... OK
Applying learning_logs.0002_entry... OK
Applying learning_logs.0003_topic_user... OK
Applying sessions.0001_initial... OK
(ll_env)learning_log$
```

---

We first issue the command `heroku run python manage.py migrate` ❶. Heroku then creates a terminal session to run the `migrate` command ❷. At ❸ Django applies the default migrations and the migrations we generated during the development of Learning Log.

Now when you visit your deployed app, you should be able to use it just as you did on your local system. However, you won't see any of the data you entered on your local deployment, because we didn't copy the data to the live server. This is normal practice: you don't usually copy local data to a live deployment because the local data is usually test data.

You can share your Heroku link to let anyone use your version of Learning Log. In the next section we'll complete a few more tasks to finish the deployment process and set you up to continue developing Learning Log.

## Refining the Heroku Deployment

In this section we'll refine the deployment by creating a superuser, just as we did locally. We'll also make the project more secure by changing the setting `DEBUG` to `False`, so users won't see any extra information in error messages that they could use to attack the server.

### Creating a Superuser on Heroku

You've already seen that we can run one-off commands using the `heroku run` command. But you can also run commands by opening a Bash terminal session while connected to the Heroku server using the command `heroku run bash`. *Bash* is the language that runs in many Linux terminals. We'll use the Bash terminal session to create a superuser so we can access the admin site on the live app:

---

```
(ll_env)learning_log$ heroku run bash
Running `bash` on afternoon-meadow-2775... up, run.6244
❶ ~ $ ls
learning_log learning_logs manage.py Procfile requirements.txt runtime.txt
users
staticfiles
❷ ~ $ python manage.py createsuperuser
Username (leave blank to use 'u41907'): ll_admin
Email address:
Password:
Password (again):
Superuser created successfully.
❸ ~ $ exit
exit
(ll_env)learning_log$
```

---

At ❶ we run `ls` to see which files and directories exist on the server, which should be the same files we have on our local system. You can navigate this file system like any other.

#### NOTE

*Windows users will use the same commands shown here (such as `ls` instead of `dir`), because you're running a Linux terminal through a remote connection.*

At ❷ we run the command to create a superuser, which outputs the same prompts we saw on our local system when we created a superuser in Chapter 18. When you're finished creating the superuser in this terminal session, use the `exit` command to return to your local system's terminal session ❸.

Now you can add `/admin/` to the end of the URL for the live app and log in to the admin site. For me, the URL is `https://afternoon-meadow-2775.herokuapp.com/admin/`.

If other people have already started using your project, be aware that you'll have access to all of their data! Don't take this lightly, and users will continue to trust you with their data.

## Creating a User-Friendly URL on Heroku

You'll probably want your URL to be friendlier and more memorable than `https://afternoon-meadow-2775.herokuapp.com/`. You can rename the app using a single command:

---

```
(ll_env)learning_log$ heroku apps:rename learning-log
Renaming afternoon-meadow-2775 to learning-log... done
https://learning-log.herokuapp.com/ | https://git.heroku.com/learning-log.git
Git remote heroku updated
(ll_env)learning_log$
```

---

You can use letters, numbers, and dashes when naming your app, and call it whatever you want, as long as no one else has claimed the name. This deployment now lives at `https://learning-log.herokuapp.com/`. The project is no longer available at the previous URL; the `apps:rename` command completely moves the project to the new URL.

### NOTE

*When you deploy your project using Heroku's free service, Heroku puts your deployment to sleep if it hasn't received any requests after a certain amount of time or if it's been too active for the free tier. The first time a user accesses the site after it's been sleeping, it will take longer to load, but the server will respond to subsequent requests more quickly. This is how Heroku can afford to offer free deployments.*

## Securing the Live Project

One glaring security issue exists in the way our project is currently deployed: the setting `DEBUG=True` in `settings.py`, which provides debug messages when errors occur. Django's error pages give you vital debugging information when you're developing a project, but they give way too much information to attackers if you leave them enabled on a live server. We also need to make sure no one can get information or redirect requests by pretending to be the project's host.

Let's modify `settings.py` so we can see error messages locally but not on the live deployment:

---

```
settings.py  --snip--
# Heroku settings
if os.getcwd() == '/app':
    --snip--
    # Honor the 'X-Forwarded-Proto' header for request.is_secure().
    SECURE_PROXY_SSL_HEADER = ('HTTP_X_FORWARDED_PROTO', 'https')

    # Allow only Heroku to host the project.
    ❶ ALLOWED_HOSTS = ['learning-log.herokuapp.com']

    ❷ DEBUG = False

    # Static asset configuration
    --snip--
```

---

We need to make only two changes: at ❶ we modify `ALLOWED_HOSTS`, so the only server allowed to host the project is Heroku. You need to use the name of your app, whether it's the name Heroku provided, such as *afternoon-meadow-2775.herokuapp.com*, or the name you chose. At ❷ we set `DEBUG` to `False`, so Django won't share sensitive information when an error occurs.

## Committing and Pushing Changes

Now we need to commit the changes made to *settings.py* to the Git repository, and then push the changes to Heroku. Here's a terminal session showing this process:

---

```
❶ (ll_env)learning_log$ git commit -am "Set DEBUG=False for Heroku."
[master 081f635] Set DEBUG=False for Heroku.
 1 file changed, 4 insertions(+), 2 deletions(-)
❷ (ll_env)learning_log$ git status
# On branch master
nothing to commit, working directory clean
(ll_env)learning_log$
```

---

We issue the `git commit` command with a short but descriptive commit message ❶. Remember that the `-am` flag makes sure Git commits all the files that have changed and records the log message. Git recognizes that one file has changed and commits this change to the repository.

At ❷ the status shows that we're working on the master branch of the repository and that there are now no new changes to commit. It's essential that you check the status for this message before pushing to Heroku. If you don't see this message, some changes haven't been committed, and those changes won't be pushed to the server. You can try issuing the `commit` command again, but if you're not sure how to resolve the issue, read through Appendix D to better understand how to work with Git.

Now let's push the updated repository to Heroku:

---

```
(ll_env)learning_log$ git push heroku master
--snip--
remote: -----> Python app detected
remote: -----> Installing dependencies with pip
--snip--
remote: -----> Launching... done, v8
remote:      https://learning-log.herokuapp.com/ deployed to Heroku
remote: Verifying deploy.... done.
To https://git.heroku.com/learning-log.git
  4c9d111..ef65d2b master -> master
(ll_env)learning_log$
```

---

Heroku recognizes that the repository has been updated, and it rebuilds the project to make sure all the changes have been taken into account. It doesn't rebuild the database, so we won't have to run `migrate` for this update.

To check that the deployment is more secure now, enter the URL of your project with an extension we haven't defined. For example, try to visit <http://learning-log.herokuapp.com/letmein/>. You should see a generic error page on your live deployment that doesn't give away any specific information about the project. If you try the same request on the local version of Learning Log at <http://localhost:8000/letmein/>, you should see the full Django error page. The result is perfect: you'll see informative error messages when you're developing the project further, but users won't see critical information about the project's code.

## Creating Custom Error Pages

In Chapter 19, we configured Learning Log to return a 404 error if the user requests a topic or entry that doesn't belong to them. You've probably seen some 500 server errors (internal errors) by this point as well. A 404 error usually means your Django code is correct, but the object being requested doesn't exist; a 500 error usually means there's an error in the code you've written, such as an error in a function in *views.py*. Currently, Django returns the same generic error page in both situations, but we can write our own 404 and 500 error page templates that match the overall appearance of Learning Log. These templates must go in the root template directory.

## Making Custom Templates

In the *learning\_log/learning\_log* folder, make a new folder called *templates*. Then make a new file called *404.html* using the following code:

---

```
404.html  {% extends "learning_logs/base.html" %}

          {% block header %}
            <h2>The item you requested is not available. (404)</h2>
          {% endblock header %}
```

---

This simple template provides the generic 404 error page information but is styled to match the rest of the site.

Make another file called *500.html* using the following code:

---

```
500.html  {% extends "learning_logs/base.html" %}

          {% block header %}
            <h2>There has been an internal error. (500)</h2>
          {% endblock header %}
```

---

These new files require a slight change to *settings.py*.

---

```
settings.py  --snip--
            TEMPLATES = [
                {
                    'BACKEND': 'django.template.backends.django.DjangoTemplates',
                    'DIRS': [os.path.join(BASE_DIR, 'learning_log/templates')],
```

```
        'APP_DIRS': True,
        --snip--
    },
]
--snip--
```

---

This change tells Django to look in the root template directory for the error page templates.

### Viewing the Error Pages Locally

If you want to see what the error pages look like on your system before pushing them to Heroku, you'll first need to set `Debug=False` on your local settings to suppress the default Django debug pages. To do so, make the following changes to `settings.py` (make sure you're working in the part of `settings.py` that applies to the local environment, not the part that applies to Heroku):

```
settings.py  --snip--
# SECURITY WARNING: don't run with debug turned on in production!
DEBUG = False

ALLOWED_HOSTS = ['localhost']
--snip--
```

---

You must have at least one host specified in `ALLOWED_HOSTS` when `DEBUG` is set to `False`. Now request a topic or entry that doesn't belong to you to see the 404 error page, and request a URL that doesn't exist (such as `localhost:8000/letmein/`) to see the 500 error page.

When you're finished checking the error pages, set `DEBUG` back to `True` to further develop Learning Log. (Make sure `DEBUG` is still set to `False` in the section of `settings.py` that applies to the Heroku deployment.)

#### NOTE

*The 500 error page won't show any information about the user who's logged in, because Django doesn't send any context information in the response when there's a server error.*

### Pushing the Changes to Heroku

Now we need to commit the template changes and push them live to Heroku:

- 
- ```
❶ (ll_env)learning_log$ git add .
❷ (ll_env)learning_log$ git commit -am "Added custom 404 and 500 error pages."
   3 files changed, 15 insertions(+), 10 deletions(-)
   create mode 100644 learning_log/templates/404.html
   create mode 100644 learning_log/templates/500.html
❸ (ll_env)learning_log$ git push heroku master
--snip--
remote: Verifying deploy.... done.
```

```
To https://git.heroku.com/learning-log.git
2b34ca1..a64d8d3 master -> master
(ll_env)learning_log$
```

---

We issue the `git add .` command at ❶ because we created some new files in the project, so we need to tell Git to start tracking these files. Then we commit the changes ❷ and push the updated project to Heroku ❸.

Now when an error page appears, it should have the same styling as the rest of the site, making for a smoother user experience when errors arise.

### Using the `get_object_or_404()` Method

At this point, if a user manually requests a topic or entry that doesn't exist, they'll get a 500 server error. Django tries to render the page but it doesn't have enough information to do so, and the result is a 500 error. This situation is more accurately handled as a 404 error, and we can implement this behavior with the Django shortcut function `get_object_or_404()`. This function tries to get the requested object from the database, but if that object doesn't exist, it raises a 404 exception. We'll import this function into `views.py` and use it in place of `get()`:

---

```
views.py --snip--
from django.shortcuts import render, get_object_or_404
from django.http import HttpResponseRedirect, Http404
--snip--
@login_required
def topic(request, topic_id):
    """Show a single topic and all its entries."""
    topic = get_object_or_404(Topic, id=topic_id)
    # Make sure the topic belongs to the current user.
    --snip--
```

---

Now when you request a topic that doesn't exist (for example, `http://localhost:8000/topics/999999/`), you'll see a 404 error page. To deploy this change, make a new commit, and then push the project to Heroku.

### Ongoing Development

You might want to further develop Learning Log after your initial push to a live server or develop your own projects to deploy. There's a fairly consistent process for updating projects.

First, you'll make any changes needed to your local project. If your changes result in any new files, add those files to the Git repository using the command `git add .` (be sure to include the dot at the end of the command). Any change that requires a database migration will need this command, because each migration generates a new migration file.

Then commit the changes to your repository using `git commit -am "commit message"`. Thereafter, push your changes to Heroku using the command `git push heroku master`. If you migrated your database locally, you'll need to migrate the live database as well. You can either use the one-off

command `heroku run python manage.py migrate`, or open a remote terminal session with `heroku run bash` and run the command `python manage.py migrate`. Then visit your live project, and make sure the changes you expect to see have taken effect.

It's easy to make mistakes during this process, so don't be surprised when something goes wrong. If the code doesn't work, review what you've done and try to spot the mistake. If you can't find the mistake or you can't figure out how to undo the mistake, refer to the suggestions for getting help in Appendix C. Don't be shy about asking for help: everyone else learned to build projects by asking the same questions you're likely to ask, so someone will be happy to help you. Solving each problem that arises helps you steadily develop your skills until you're building meaningful, reliable projects and you're answering other people's questions as well.

## ***The SECRET\_KEY Setting***

Django uses the value of the `SECRET_KEY` setting in `settings.py` to implement a number of security protocols. In this project, we've committed our settings file to the repository with the `SECRET_KEY` setting included. This is fine for a practice project, but the `SECRET_KEY` setting should be handled more carefully for a production site. If you build a project that's getting meaningful use, make sure you research how to handle your `SECRET_KEY` setting more securely.

## ***Deleting a Project on Heroku***

It's great practice to run through the deployment process a number of times with the same project or with a series of small projects to get the hang of deployment. But you'll need to know how to delete a project that's been deployed. Heroku might also limit the number of projects you can host for free, and you don't want to clutter your account with practice projects.

Log in to the Heroku website (<https://heroku.com/>), and you'll be redirected to a page showing a list of your projects. Click the project you want to delete, and you'll see a new page with information about the project. Click the **Settings** link, and scroll down until you see a link to delete the project. This action can't be reversed, so Heroku will ask you to confirm the request for deletion by manually entering the project's name.

If you prefer working from a terminal, you can also delete a project by issuing the `destroy` command:

---

```
(ll_env)learning_log$ heroku apps:destroy --app appname
```

---

Here *appname* is the name of your project, which is either something like `afternoon-meadow-2775` or `learning-log` if you've renamed the project. You'll be prompted to reenter the project name to confirm the deletion.

### **NOTE**

*Deleting a project on Heroku does nothing to your local version of the project. If no one has used your deployed project and you're just practicing the deployment process, it's perfectly reasonable to delete your project on Heroku and redeploy it.*



### TRY IT YOURSELF

**20-3. Live Blog:** Deploy the Blog project you’ve been working on to Heroku. Make sure you set `DEBUG` to `False` and change the `ALLOWED_HOSTS` setting, so your deployment is reasonably secure.

**20-4. More 404s:** The `get_object_or_404()` function should also be used in the `new_entry()` and `edit_entry()` views. Make this change, test it by entering a URL like `http://localhost:8000/new_entry/99999/`, and check that you see a 404 error.

**20-5. Extended Learning Log:** Add one feature to Learning Log, and push the change to your live deployment. Try a simple change, such as writing more about the project on the home page. Then try adding a more advanced feature, such as giving users the option of making a topic public. This would require an attribute called `public` as part of the `Topic` model (this should be set to `False` by default) and a form element on the `new_topic` page that allows the user to change a topic from private to public. You’d then need to migrate the project and revise `views.py` so any topic that’s public is visible to unauthenticated users as well. Remember to migrate the live database after you’ve pushed your changes to Heroku.

## Summary


In this chapter you learned to give your projects a simple but professional appearance using the Bootstrap library and the `django-bootstrap3` app. Using Bootstrap means the styles you choose will work consistently on almost any device people use to access your project.

You learned about Bootstrap’s templates, and we used the *Static top navbar* template to create a simple look and feel for Learning Log. You learned how to use a jumbotron to make a home page’s message stand out, and you learned to style all the pages in a site consistently.

In the final part of the project, you learned how to deploy a project to Heroku’s servers so anyone can access it. You made a Heroku account and installed some tools that help manage the deployment process. You used Git to commit the working project to a repository and then pushed the repository to Heroku’s servers. Finally, you learned to secure your app by setting `DEBUG=False` on the live server.

Now that you’ve finished Learning Log, you can start building your own projects. Start simple, and make sure the project works before adding complexity. Enjoy your learning, and good luck with your projects!

## AFTERWORD



Congratulations! You've learned the basics of Python and applied your knowledge to meaningful projects. You've made a game, visualized some data, and made a web application. From here, you can go in a number of different directions to continue developing your programming skills.

First, you should continue to work on meaningful projects that interest you. Programming is more appealing when you're solving relevant and significant problems, and you now have the skills to engage in a variety of projects. You could invent your own game or write your own version of a classic arcade game. You might want to explore some data that's important to you and make visualizations that show interesting patterns and connections. You could create your own web application or try to emulate one of your favorite apps.

Whenever possible, invite other people to try using your programs. If you write a game, let other people play it. If you make a visualization, show it to others and see if it makes sense to them. If you make a web app, deploy it online and invite others to try it out. Listen to your users and try to incorporate their feedback into your projects; you'll become a better programmer if you do.


When you work on your own projects, you'll run into problems that are challenging, or even impossible, to solve on your own. Keep finding ways to ask for help, and find your own place in the Python community. Join a local Python User Group or explore some online Python communities. Consider attending a PyCon near you as well.

You should strive to maintain a balance between working on projects that interest you and developing your Python skills in general. Many Python learning sources are available online, and a large number of Python books target intermediate programmers. Many of these resources will be accessible to you now that you know the basics and how to apply your skills. Working through Python tutorials and books will build directly on what you learned here and deepen your understanding of programming in general and Python in particular. Then when you go back to working on projects after focusing on learning about Python, you'll be capable of solving a wider variety of problems more efficiently.

Congratulations on how far you've come, and good luck with your continued learning!

# D

## USING GIT FOR VERSION CONTROL



Version control software allows you to take snapshots of a project whenever it's in a working state. When you make changes to a project—for example, when you implement a new feature—you have the option of reverting back to a previous working state if the project's current state isn't functioning well.

Using version control software gives you the freedom to work on improvements and make mistakes without worrying about ruining your project. This is especially critical in large projects, but can also be helpful in smaller projects, even when you're working on programs contained in a single file.

In this appendix you'll learn to install Git and use it for version control in the programs you're working on now. Git is the most popular version control software in use today. Many of its advanced tools help teams

collaborate on large projects, but its most basic features also work well for solo developers. Git implements version control by tracking the changes made to every file in a project; if you make a mistake, you can just return to a previously saved state.

## Installing Git

Git runs on all operating systems, but there are different approaches to installing it on each system. The following sections provide specific instructions for each operating system.

### *Installing Git on Linux*

To install Git on Linux, enter the following:

---

```
$ sudo apt-get install git
```

---

That's it. You can now use Git in your projects.

### *Installing Git on OS X*

Git may already be installed on your system, so try issuing the command `git --version`. If you see output listing a specific version number, Git is installed on your system. If you see a message prompting you to install or update Git, simply follow the onscreen directions.

You can also go to <https://git-scm.com/>, follow the Downloads link, and click an appropriate installer for your system.

### *Installing Git on Windows*

You can install Git for Windows from <http://msysgit.github.io/>.

### *Configuring Git*

Git keeps track of who makes changes to a project, even when there's only one person working on the project. To do this, Git needs to know your username and email. You have to provide a username, but feel free to make up a fake email address:

---

```
$ git config --global user.name "username"
$ git config --global user.email "username@example.com"
```

---

If you forget this step, Git will prompt you for this information when you make your first commit.

## Making a Project

Let's make a project to work with. Create a folder somewhere on your system called *git\_practice*. Inside the folder, make a simple Python program:

---

```
hello_world.py    print("Hello Git world!")
```

---

We'll use this program to explore Git's basic functionality.

## Ignoring Files

Files with the extension *.pyc* are automatically generated from *.py* files, so we don't need Git to keep track of them. These files are stored in a directory called *\_\_pycache\_\_*. To tell Git to ignore this directory, make a special file called *.gitignore*—with a dot at the beginning of the filename and no file extension—and add the following line to it:

---

```
.gitignore    __pycache__ /
```

---

This tells Git to ignore any file in the *\_\_pycache\_\_* directory. Using a *.gitignore* file will keep your project clutter free and easier to work with.

### NOTE

*If you're using Python 2.7, replace this line with `*.pyc`. Python 2.7 doesn't create a `__pycache__` directory; each `.pyc` file is stored in the same directory as its corresponding `.py` file. The asterisk tells Git to ignore any file with the `.pyc` extension.*

You might need to modify your text editor's settings so it will show hidden files in order to open *.gitignore*. Some editors are set to ignore filenames that begin with a dot.

## Initializing a Repository

Now that you have a directory containing a Python file and a *.gitignore* file, you can initialize a Git repository. Open a terminal, navigate to the *git\_practice* folder, and run the following command:

---

```
git_practice$ git init
Initialized empty Git repository in git_practice/.git/
git_practice$
```

---

The output shows that Git has initialized an empty repository in *git\_practice*. A *repository* is the set of files in a program that Git is actively tracking. All the files Git uses to manage the repository are located in the hidden directory *.git/*, which you won't need to work with at all. Just don't delete that directory, or you'll lose your project's history.

## Checking the Status

Before doing anything else, let's look at the status of the project:

---

```
git_practice$ git status
❶ # On branch master
#
# Initial commit
#
❷ # Untracked files:
#   (use "git add <file>..." to include in what will be committed)
#
#   .gitignore
#   hello_world.py
#
❸ nothing added to commit but untracked files present (use "git add" to track)
git_practice$
```

---

In Git, a *branch* is a version of the project you're working on; here you can see that we're on a branch named master ❶. Each time you check your project's status, it should say that you're on the branch master. We then see that we're about to make the initial commit. A *commit* is a snapshot of the project at a particular point in time.

Git informs us that untracked files are in the project ❷, because we haven't told it which files to track yet. Then we're told that there's nothing added to the current commit, but there are untracked files present that we might want to add to the repository ❸.

## Adding Files to the Repository

Let's add the two files to the repository, and check the status again:

---

```
❶ git_practice$ git add .
❷ git_practice$ git status
# On branch master
#
# Initial commit
#
# Changes to be committed:
#   (use "git rm --cached <file>..." to unstage)
#
❸ # new file:   .gitignore
# new file:   hello_world.py
#
git_practice$
```

---

The command `git add .` adds all files within a project that are not already being tracked to the repository ❶. It doesn't commit the files; it just tells Git to start paying attention to them. When we check the status of the

project now, we can see that Git recognizes some changes that need to be committed ❷. The label *new file* means these files were newly added to the repository ❸.

## Making a Commit

Let's make the first commit:

---

```
❶ git_practice$ git commit -m "Started project."
❷ [master (root-commit) c03d2a3] Started project.
❸ 2 files changed, 1 insertion(+)
   create mode 100644 .gitignore
   create mode 100644 hello_world.py
❹ git_practice$ git status
# On branch master
nothing to commit, working directory clean
git_practice$
```

---

We issue the command `git commit -m "message"` ❶ to take a snapshot of the project. The `-m` flag tells Git to record the message that follows ("Started project.") in the project's log. The output shows that we're on the master branch ❷ and that two files have changed ❸.

When we check the status now, we can see that we're on the master branch, and we have a clean working directory ❹. This is the message you want to see each time you commit a working state of your project. If you get a different message, read it carefully; it's likely you forgot to add a file before making a commit.

## Checking the Log

Git keeps a log of all commits made to the project. Let's check the log:

---

```
git_practice$ git log
commit a9d74d87f1aa3b8f5b2688cb586eac1a908cfc7f
Author: Eric Matthes <eric@example.com>
Date:   Mon Mar 16 07:23:32 2015 -0800
```

```
Started project.
git_practice$
```

---

Each time you make a commit, Git generates a unique, 40-character reference ID. It records who made the commit, when it was made, and the message recorded. You won't always need all of this information, so Git provides an option to print a simpler version of the log entries:

---

```
git_practice$ git log --pretty=oneline
a9d74d87f1aa3b8f5b2688cb586eac1a908cfc7f Started project.
git_practice$
```

---



The `--pretty=oneline` flag provides the two most important pieces of information: the reference ID of the commit and the message recorded for the commit.

## The Second Commit

To see the real power of version control, we need to make a change to the project and commit that change. Here we'll just add another line to *hello\_world.py*:

```
hello_world.py print("Hello Git world!")
                print("Hello everyone.")
```

If we check the status of the project, we'll see that Git has noticed the file that changed:

```
git_practice$ git status
❶ # On branch master
# Changes not staged for commit:
#   (use "git add <file>..." to update what will be committed)
#   (use "git checkout -- <file>..." to discard changes in working directory)
#
❷ #    modified:   hello_world.py
#
❸ no changes added to commit (use "git add" and/or "git commit -a")
git_practice$
```

We see the branch we're working on **❶**, the name of the file that was modified **❷**, and that no changes have been committed **❸**. Let's commit the change and check the status again:

```
❶ git_practice$ git commit -am "Extended greeting."
[master 08d4d5e] Extended greeting.
 1 file changed, 1 insertion(+)
❷ git_practice$ git status
# On branch master
nothing to commit, working directory clean
❸ git_practice$ git log --pretty=oneline
08d4d5e39cb906f6cff197bd48e9ab32203d7ed6 Extended greeting.
be017b7f06d390261dbc64ff593be6803fd2e3a1 Started project.
git_practice$
```

We make a new commit, passing the `-am` flag when we use the command `git commit` **❶**. The `-a` flag tells Git to add all modified files in the repository to the current commit. (If you create any new files between commits, simply reissue the `git add .` command to include the new files in the repository.) The `-m` flag tells Git to record a message in the log for this commit.

When we check the status of the project, we see that we once again have a clean working directory **❷**. Finally, we see the two commits in the log **❸**.

## Reverting a Change

Now let's see how to abandon a change and revert back to the previous working state. First, add a new line to *hello\_world.py*:

```
hello_world.py print("Hello Git world!")
                print("Hello everyone.")

                print("Oh no, I broke the project!")
```

Save and run this file.

We check the status and see that Git notices this change:

```
git_practice$ git status
# On branch master
# Changes not staged for commit:
#   (use "git add <file>..." to update what will be committed)
#   (use "git checkout -- <file>..." to discard changes in working directory)
#
❶ #    modified:   hello_world.py
#
no changes added to commit (use "git add" and/or "git commit -a")
git_practice$
```

Git sees that we modified *hello\_world.py* ❶, and we can commit the change if we want to. But this time, instead of committing the change, we want to revert back to the last commit when we knew our project was working. We won't do anything to *hello\_world.py*; we won't delete the line or use the Undo feature in the text editor. Instead, enter the following commands in your terminal session:

```
git_practice$ git checkout .
git_practice$ git status
# On branch master
nothing to commit, working directory clean
git_practice$
```

The command `git checkout .` allows you to work with any previous commit. The command `git checkout .` abandons any changes made since the last commit and restores the project to the last committed state.

When you return to your text editor, you'll see that *hello\_world.py* has changed back to this:

```
print("Hello Git world!")
print("Hello everyone.")
```

Although going back to a previous state may seem trivial in this simple project, if we were working on a large project with dozens of modified files, all of the files that had changed since the last commit would be reverted. This feature is incredibly useful: you can make as many changes as you

want when implementing a new feature, and if they don't work, you can discard them without harming the project. You don't have to remember those changes and manually undo them. Git does all of that for you.

**NOTE**

*You might have to click in your editor's window to refresh the file and see the previous version.*

## Checking Out Previous Commits

You can check out any commit in your log, not just the most recent, by including the first six characters of the reference ID instead of a dot. By checking it out, you can review an earlier commit, and you're able to then return to the latest commit or abandon your recent work and pick up development from the earlier commit:

---

```
git_practice$ git log --pretty=oneline
08d4d5e39cb906f6cfff197bd48e9ab32203d7ed6 Extended greeting.
be017b7f06d390261dbc64ff593be6803fd2e3a1 Started project.
git_practice$ git checkout be017b
Note: checking out 'be017b'.
```

- ❶ You are in 'detached HEAD' state. You can look around, make experimental changes and commit them, and you can discard any commits you make in this state without impacting any branches by performing another checkout.

If you want to create a new branch to retain commits you create, you may do so (now or later) by using `-b` with the checkout command again. Example:

```
git checkout -b new_branch_name

HEAD is now at be017b7... Started project.
git_practice$
```

---

When you check out a previous commit, you leave the master branch and enter what Git refers to as a *detached HEAD* state ❶. *HEAD* is the current state of the project; we are *detached* because we've left a named branch (master, in this case).

To get back to the master branch, you check it out:

---

```
git_practice$ git checkout master
Previous HEAD position was be017b7... Started project.
Switched to branch 'master'
git_practice$
```

---

This brings you back to the master branch. Unless you want to work with some more advanced features of Git, it's best not to make any changes to your project when you've checked out an old commit. However, if you're

the only one working on a project and you want to discard all of the more recent commits and go back to a previous state, you can reset the project to a previous commit. Working from the master branch, enter the following:

---

```
❶ git_practice$ git status
# On branch master
nothing to commit, working directory clean
❷ git_practice$ git log --pretty=oneline
08d4d5e39cb906f6cfff197bd48e9ab32203d7ed6 Extended greeting.
be017b7f06d390261dbc64ff593be6803fd2e3a1 Started project.
❸ git_practice$ git reset --hard be017b
HEAD is now at be017b7 Started project.
❹ git_practice$ git status
# On branch master
nothing to commit, working directory clean
❺ git_practice$ git log --pretty=oneline
be017b7f06d390261dbc64ff593be6803fd2e3a1 Started project.
git_practice$
```

---

We first check the status to make sure we're on the master branch ❶. When we look at the log, we see both commits ❷. We then issue the `git reset --hard` command with the first six characters of the reference ID of the commit we want to revert to permanently ❸. We check the status again and see we're on the master branch with nothing to commit ❹. When we look at the log again, we see that we're at the commit we wanted to start over from ❺.

## Deleting the Repository

Sometimes you'll mess up your repository's history and won't know how to recover it. If this happens, first consider asking for help using the methods discussed in Appendix C. If you can't fix it and you're working on a solo project, you can continue working with the files but get rid of the project's history by deleting the `.git` directory. This won't affect the current state of any of the files, but it will delete all commits, so you won't be able to check out any other states of the project.

To do this, either open a file browser and delete the `.git` repository or do it from the command line. Afterwards, you'll need to start over with a fresh repository to start tracking your changes again. Here's what this entire process looks like in a terminal session:

---

```
❶ git_practice$ git status
# On branch master
nothing to commit, working directory clean
❷ git_practice$ rm -rf .git
❸ git_practice$ git status
fatal: Not a git repository (or any of the parent directories): .git
❹ git_practice$ git init
Initialized empty Git repository in git_practice/.git/
```

```

❸ git_practice$ git status
# On branch master
#
# Initial commit
#
# Untracked files:
#   (use "git add <file>..." to include in what will be committed)
#
#   .gitignore
#   hello_world.py
#
nothing added to commit but untracked files present (use "git add" to track)
❹ git_practice$ git add .
git_practice$ git commit -m "Starting over."
[master (root-commit) 05f5e01] Starting over.
 2 files changed, 2 insertions(+)
 create mode 100644 .gitignore
 create mode 100644 hello_world.py
❺ git_practice$ git status
# On branch master
nothing to commit, working directory clean
git_practice$

```

---

We first check the status and see that we have a clean working directory ❶. Then we use the command `rm -rf .git` to delete the `.git` directory (`rmdir /s .git` on Windows) ❷. When we check the status after deleting the `.git` folder, we're told that this is not a Git repository ❸. All the information Git uses to track a repository is stored in the `.git` folder, so removing it deletes the entire repository.

We're then free to use `git init` to start a fresh repository ❹. Checking the status shows that we're back at the initial stage, awaiting the first commit ❺. We add the files and make the first commit ❻. Checking the status now shows us that we're on the new `master` branch with nothing to commit ❼.

Using version control takes a bit of practice, but once you start using it you'll never want to work without it again.

# RESOURCES

Visit <https://www.nostarch.com/pythoncrashcourse/> for resources, errata, and more information.

More no-nonsense books from



**NO STARCH PRESS**




## **PYTHON PLAYGROUND** Geeky Weekend Projects for the Curious Programmer

by MAHESH VENKITACHALAM  
OCTOBER 2015, 352 PP., \$29.95  
ISBN 978-1-59327-604-1




## **AUTOMATE THE BORING STUFF WITH PYTHON** Practical Programming for Total Beginners

by AL SWEIGART  
APRIL 2015, 504 PP., \$29.95  
ISBN 978-1-59327-599-0




## **DOING MATH WITH PYTHON** Use Programming to Explore Algebra, Statistics, Calculus, and More!

by AMIT SAHA  
AUGUST 2015, 264 PP., \$29.95  
ISBN 978-1-59327-640-9




## **THE LINUX COMMAND LINE** A Complete Introduction

by WILLIAM E. SHOTTS, JR.  
JANUARY 2012, 480 PP., \$39.95  
ISBN 978-1-59327-389-7



## **ELOQUENT JAVASCRIPT, 2ND EDITION** A Modern Introduction to Programming

by MARIJN HAVERBEKE  
DECEMBER 2014, 472 PP., \$39.95  
ISBN 978-1-59327-584-6



## **BLACK HAT PYTHON** Python Programming for Hackers and Pentesters

by JUSTIN SEITZ  
DECEMBER 2014, 192 PP., \$34.95  
ISBN 978-1-59327-590-7

### PHONE:

800.420.7240 OR  
415.863.9900


### EMAIL:

[SALES@NOSTARCH.COM](mailto:SALES@NOSTARCH.COM)

### WEB:

[WWW.NOSTARCH.COM](http://WWW.NOSTARCH.COM)

## LEARN PYTHON— FAST!



*Python Crash Course* is a fast-paced, thorough introduction to programming with Python that will have you writing programs, solving problems, and making things that work in no time.

In the first half of the book, you'll learn about basic programming concepts, such as lists, dictionaries, classes, and loops, and practice writing clean and readable code with exercises for each topic. You'll also learn how to make your programs interactive and how to test your code safely before adding it to a project. In the second half of the book, you'll put your new knowledge into practice with three substantial projects: a Space Invaders–inspired arcade game, data visualizations with Python's super-handly libraries, and a simple web app you can deploy online.

As you work through *Python Crash Course*, you'll learn how to:

- Use powerful Python libraries and tools, including matplotlib, NumPy, and Pygal

- Make 2D games that respond to keypresses and mouse clicks, and that grow more difficult as the game progresses
- Work with data to generate interactive visualizations
- Create and customize simple web apps and deploy them safely online
- Deal with mistakes and errors so you can solve your own programming problems

If you've been thinking seriously about digging into programming, *Python Crash Course* will get you up to speed and have you writing real programs fast. Why wait any longer? Start your engines and code!

### ABOUT THE AUTHOR

Eric Matthes is a high school science and math teacher living in Alaska, where he teaches an introductory Python course. He has been writing programs since he was five years old.

## COVERS PYTHON 2 AND 3



THE FINEST IN GEEK ENTERTAINMENT™  
[www.nostarch.com](http://www.nostarch.com)

"I LIE FLAT."

*This book uses RepKover—a durable binding that won't snap shut*

ISBN: 978-1-59327-603-4



**\$39.95 (\$45.95 CDN)**



SHELF LIFE:  
PROGRAMMING LANGUAGES/  
PYTHON

WOW! eBook  
[www.wowebook.org](http://www.wowebook.org)